

# Lecture 3: OOP, classes, scripts and some other issues

- Everything is an object, so... everything!
- OOP
- Classes
- Some brief notes on other language features
- Sets and choice of data structures
- Scripting
- Presentation by Anders Mäarak Leffler
- Attribution: slightly extends work by Johan Falkenjack.
- License: [CC-BY-SA 4.0 \(https://creativecommons.org/licenses/by-sa/4.0/\)](https://creativecommons.org/licenses/by-sa/4.0/)

## OOP in principle

- Paradigm. Structuring code.
  - Abstractions.
  - Rough resemblance between objects in the world, and the code.
  - Keeping related data and behaviour together.
    - Roughly: a bunch of data and the way they interact with the world.

## Objects have state (and groups related data together)

In [ ]:

```
# Most extreme (somewhat strawmannish) contrast: entirely ungrouped values.
car_1_maker = "Volvo"
car_1_model = 240
car_1_colour = "black"
car_1_cost = 90000

car_2_maker = "Toyota"
car_2_model = "Camry"
car_2_colour = "red"
car_2_cost = 100000

# ...

# How do we pass a list of cars to a function which calculates eg if Toyotas are more expensive?
```

In [3]:

```
# Grouping data together.
```

```
class Car:  
    pass    # More Later!
```

```
car_1 = Car()  
car_1.maker = "Volvo"  
car_1.model = 240  
car_1.colour = "black"  
car_1.cost = 90000
```

```
car_2 = Car()  
car_2.maker = "Volvo"  
car_2.model = "Camry"  
car_2.colour = "red"  
car_2.cost = 10000
```

```
cars = [car_1, car_2]  
cars  
cars[0].colour  
car_1.cost = 99999999999999  
cars[0].cost
```

Out[3]:

99999999999999

Note: not only grouping, but now we can change the cost of a *specific* car.

In [4]:

```
# Grouping data together, somewhat more structured fashion.
# Homespun version to avoid the code repetition above. (not part of
# Python)
def bad_init(maker, model, cost, colour):
    res = Car()
    res.maker = maker
    res.model = model
    res.cost = cost
    res.colour = colour
    return res

car_test = bad_init("volvo", 240, 12313, "white")

# Creating a "blueprint" for what a car (in our application) _is_.

class Car:
    # Somehow defining that a car _is_ something which has all these properties.
    # (in our application)

    def __init__(self, maker, model, cost, colour = "black"):
        self.maker = maker
        self.model = model
        self.cost = cost
        self.colour = colour

car_1 = Car(maker = "Volvo", model = 240, cost = 90000) # black is the default colour
car_2 = Car(maker = "Toyota", model = "Camry", colour = "red", cost = 10000)
cars = [car_1, car_2]
cars
```

Out[4]:

```
[<__main__.Car at 0x7f32fd139e10>, <__main__.Car at 0x7f32fd139f28>]
```

In [10]:

```
cars = [Car(maker = "Volvo", model = 240, colour = "black", cost = 90000),
        Car(maker = "Toyota", model = "Camry", colour = "red", cost = 10000)]
cars
```

Out[10]:

```
[<__main__.Car at 0x7fc1f80d3e80>, <__main__.Car at 0x7fc1f80d3e10>]
```

In [ ]:

```
# With a data source such as an API connection which gives us data in a certain format,
a CSV file handle,
# we can automate this.

# Nothing specific to OOP so far, but noteworthy feature.

some_magic_data_iterable = [("Volvo", "black", 240, 90000), ("Toyota", "red", "Camry",
10000)]

cars = [Car(maker = maker, colour = colour, model = model, cost = cost)
        for maker, colour, model, cost in some_magic_data_iterable]
```

Notes:

- Couldn't we do this via some other abstraction (namedtuple, or tuples + a function which picks out the right parts, such as `get_maker(car) == car[0]` )...?
- Now we have objects with state. What about "behaviours"?
  - (Philosophical question for those so inclined: structs vs objects.)

## Objects act on messages (and the concept of interface)

- (Public) interface: What messages does an object accept, and what does it do or return?

In [13]:

```
my_seq = [1,2,3]    # State before: my_seq is a list with these three elements.
other_seq = [4,5]
my_seq.append(999)  # What does this mean?
# Intuitive here.
# Formally: tell my_seq to change itself.
help(my_seq.append)
```

Help on built-in function append:

```
append(...) method of builtins.list instance
  L.append(object) -> None -- append object to end
```

- Objects carry their own behaviours "wrapped in" [not technically].
  - Ex: `classifier_1.classify(image)` uses `classifier_1`'s `classify` method.  
`classifier_2.classify(image)` might use some entirely different procedure.

(Contrast: some external function `classify(classifier_1, image)` , where the behaviour is not carried with the object).

## Encapsulation. Objects should be isolated, and hide implementation

- Carry their own data, or references to where to get it.
- Carry their own behaviours.
  - (Corollary) Avoids dependence on other objects' implementation.

Corollary: other parts of the program shouldn't need to know a lot about *how* your object does things. And if you change how it does things, their code shouldn't break.

## With great power comes great responsibility

- Default: Python will let you.

# Classes in Python

- We use `class` to create classes.
  - `class Cat(): ...` which defines what it means to be a `Cat` object).
- We call the classes to create *instances*. The `__init__` method is called.
  - `Cat(name = "Alonzo")` to create a cat instance. This is just a value, a single cat.
  - Usually bind this to save value for later ( `alonzo = Cat(name="Alonzo")` ).

Design note: a particular cat is an instance, what is *common to all cats* belongs in the class.

## Creating simple classes and instances

In [17]:

```
# Defining what is common to all cats.

class Cat:
    # More data goes here!

    # Initialiser
    def __init__(self, name = "Shere Khan"):
        self.name = name
        self.description = "adorable"

    def greet(self):
        return "Hi, my name is " + self.name

    # and it should have a "greet" method that returns a greeting string.

# Creating some instances.
alonzo = Cat(name = "Alonzo")
zeno = Cat(name = "Zeno")

# Calling a method (function attribute).

alonzo.greet()
#zeno.greet()
```

Out[17]:

```
'Hi, my name is Alonzo'
```

- Classes as namespaces. What attributes does `alonzo` have above?

In [18]:

```
dir(alonzo)
```

Out[18]:

```
['_class__',
 '__delattr__',
 '__dict__',
 '__dir__',
 '__doc__',
 '__eq__',
 '__format__',
 '__ge__',
 '__getattr__',
 '__gt__',
 '__hash__',
 '__init__',
 '__init_subclass__',
 '__le__',
 '__lt__',
 '__module__',
 '__ne__',
 '__new__',
 '__reduce__',
 '__reduce_ex__',
 '__repr__',
 '__setattr__',
 '__sizeof__',
 '__str__',
 '__subclasshook__',
 '__weakref__',
 'description',
 'greet',
 'name']
```

In [19]:

```
help(dir)
```

Help on built-in function dir in module builtins:

```
dir(...)
dir([object]) -> list of strings
```

If called without an argument, return the names in the current scope.  
 Else, return an alphabetized list of names comprising (some of) the attributes  
 of the given object, and of attributes reachable from it.  
 If the object supplies a method named `__dir__`, it will be used; otherwise  
 the default `dir()` logic is used and returns:  
   for a module object: the module's attributes.  
   for a class object: its attributes, and recursively the attributes  
     of its bases.  
   for any other object: its attributes, its class's attributes, and  
     recursively the attributes of its class's base classes.

- What is the `self` ? Why do we pass it along?

In [20]:

```
alonzo.greet()  
# corresponding to...  
Cat.greet(alonzo) # "self" is the alonzo object
```

Out[20]:

```
'Hi, my name is Alonzo'
```

In [28]:

```
# Added Live after question about generating objects.  
  
myobj = type("mynewtype", (object,), { "meth" : id}) # some function  
myobj.meth(2)
```

Out[28]:

```
10910432
```

In [22]:

```
# Added Live after question about adding functions.  
  
Cat.meow = lambda self: "Meow"  
alonzo.meow()
```

Out[22]:

```
'Meow'
```

In [23]:

```
alonzo.__dict__ # What's in the instance? Is there a meow?
```

Out[23]:

```
{'name': 'Alonzo', 'description': 'adorable'}
```

[C++? Think of it somewhat like a this pointer.]

- Can we access and check for the presence of attributes in other ways than trying with `.-access` as above?

In [29]:

```
help(getattr)
```

Help on built-in function getattr in module builtins:

```
getattr(...)
    getattr(object, name[, default]) -> value
```

Get a named attribute from an object; getattr(x, 'y') is equivalent to x.y.

When a default argument is given, it is returned when the attribute doesn't exist; without it, an exception is raised in that case.

In [ ]:

```
help(setattr)
```

In [ ]:

```
help(hasattr)
```

[More to come.]

- We can generate "more of the same". A note on type .

In [31]:

```
type(alonzo) # gives us the class/constructor!
```

Out[31]:

```
<__main__.Cat at 0x7fc1f808d8d0>
```

In [32]:

```
# Accessing the class of an object in a different way.
alonzo.__class__
```

Out[32]:

```
__main__.Cat
```

In [35]:

```
# Using it.
CatConstructor = type(alonzo)
fcats = CatConstructor("Forex")
fcats.name
```

Out[35]:

```
'Forex'
```

- Bonus: we can generate classes as well.



In [ ]:

```
help(type)
```

In [38]:

```
# Bonus: creating using type.
```

```
CC = type("CoolCat", (Cat,), { "a" : 5, "zoo" : lambda self : self.a })  
CC  
tomomalley = CC()  
tomomalley.zoo()
```

Out[38]:

5

(Can be used to create types on the fly.)

- Can we have several initialisers (like several constructors in C++)?

In [ ]:

```
# Not possible to have this type of polymorph. in Python!  
def __init__(self, maker, model, ...):  
    pass  
  
def __init__(self, readymadecar):  
    pass  
  
# Only the last definition survives!
```

- What about destructors? (For those used to C++).
  - Main takeaway: **Garbage collected language! No "need" for delete, delete[], free,... due to memory..**
  - But sometimes other kinds of cleanup is useful.

In [40]:

```
# Slightly contrived example.
# Check in detail on your own time.

class DataSource:
    def __init__(self, addr = None):
        self.api_connector = APIConnection(addr) # Connect to the data source via the
internet.
        # ...

    def sign_off(self):
        """Perform an ordered sign-off from the data server."""
        self.api_connector.sign_off()
        self.api_connector = None

    def __del__(self):
        if self.api_connector is not None:
            print("Someone should've sent a proper signoff ")
            self.api_connector.signal_abnormal_use() # "normal" deletion of the connector might send other signals

            # Note that api_connector:s destructor will normally be run when the GC passes
by, and
            # there are no references to it. So you should really think about who's responsible
            # for an object, if you are to do this! Also, it might not be needed.

            print("The DataSource object destructor has finished. The GC will clean up the api_connector...")

class APIConnection:
    """Dummy class."""
    def __init__(self, address):
        pass

    def sign_off(self):
        pass

    def signal_abnormal_use(self):
        pass

def f():
    print("Entered f.")
    src = DataSource(addr = "api.scb.se")
    #src.sign_off()
    print("Exiting f now. src should go out of scope and be deleted soon:ish.")
f()
```

Entered f.

Exiting f now. src should go out of scope and be deleted soon:ish.

Someone should've sent a proper signoff

The DataSource object destructor has finished. The GC will clean up the api\_connector...

## Concept: composition (has-a)

- Adding capabilities by generating objects of their own.
- Above: every `DataSource` instance has an `APIConnection` instance of its own.
  - Design choice: don't make giant `DataSource` class which has all the `APIConnection` methods inside.
  - Delegating that part of the work to a specialised class.
- See Lutz ch 31.

## Concept: Inheritance (is-a)

- Conceptually: *every X is also a Y* (and can do everything it can).
  - Adding capabilities.
  - Specialisation.
- Operationally: add properties to the objects at the right level.
- Issue: how do we handle conflicts?

In [1]:

```
class PrimordialSoupObject:
    pass

class Animal(PrimordialSoupObject):
    def __init__(self):
        self.description = "Hello, this is Dog."
        self.x = "animal"

    def make_sound(self):
        print("Running Animal's make_sound.")
        print("Grrrr.")
        print("I believe that description is:", self.description)

class SuperHero(PrimordialSoupObject):
    def __init__(self, superpower = "Flight"):
        self.description = "Na"*50 + "Batman"
        self.x = "animal"
        self.superpower = superpower

    def make_sound(self):
        print("Running SuperHero's make_sound.")
        print("Truth, justice and the american way")
        print("I believe that description is:", self.description)

    def fight(self):
        print("POW!")

class Cat(Animal, SuperHero):
    def __init__(self, name = "Shere Khan"):

        # Run parent classes' inits (early on)
        SuperHero.__init__(self) # added later
        Animal.__init__(self)

        self.description = "cat."
        self.name = name
        self.x = "animal"

    def sdfsd fsdfsd make_sound(self):
        print("Running Cat's make_sound.")
        print("Meow.")
        print("I believe that description is:", self.description)

# What will this print? What will it yield? Why?
alonzo = Cat(name = "Alonzo")
alonzo.make_sound() # comes from Animal (since the name change)
alonzo.fight() # only in SuperHero
```

Running Animal's make\_sound.  
 Grrrr.  
 I believe that description is: cat.  
 POW!

...

- Method resolution order: [\[reference \(https://www.python.org/download/releases/2.3/mro/\)\]](https://www.python.org/download/releases/2.3/mro/)

In [50]:

```
alonzo.__class__.__mro__    # In which order will I look for attributes?
```

Out[50]:

```
(__main__.Cat,  
 __main__.SuperHero,  
 __main__.Animal,  
 __main__.PrimordialSoupObject,  
 object)
```

In [48]:

```
import inspect  
inspect.getmro(Cat)
```

Out[48]:

```
(__main__.Cat,  
 __main__.Animal,  
 __main__.SuperHero,  
 __main__.PrimordialSoupObject,  
 object)
```

- Initialisers

In [55]:

```
# Trying to access the superpower (see SuperHero initialiser).  
# Crashes unless we run the SuperHero.__init__ somehow.  
alonzo.superpower
```

Out[55]:

```
'Flight'
```

In [58]:

```
alonzo.make_sound()
```

Running SuperHero's make\_sound.  
Truth, justice and the american way  
I believe that description is: cat.

In [ ]:

```
# Attempt I (super)  
# super().__init__(self) in the initialiser.
```

In [51]:

```
# Attempt II (explicit)  
# Animal.__init__ etc
```

Why care? Specialisation!

## Finding out what kind of objects we're working with

Test if it is a direct instance.

In [60]:

```
# Can we test if alonzo is a Cat?
type(alonzo) is Cat
```

Out[60]:

True

Says something about the hierarchy.

In [62]:

```
# Can we test if alonzo is an Animal?
isinstance(alonzo, Animal) # Also check the inheritances!
```

Out[62]:

True

## Conventions for "hiding" data



Dependencies, in [xkcd 1172](https://xkcd.com/1172/) (<https://xkcd.com/1172/>).

- Many languages have public/private member distinctions.
- Pessimistic note above. ([Hyrum's law \(http://www.hyrumslaw.com/\)](http://www.hyrumslaw.com/)).
- ...but let's try to hide things.
- By default everything in Python is public.

In [66]:

```
class SecretKeeper():
    def __init__(self, hidden):
        # self.hidden = hidden    # Initial code.
        self.__hidden = hidden    # Mangle name to "~hide"

c1 = SecretKeeper(hidden = "supersecret")
c1.hidden
```

```
-----
-
AttributeError                                Traceback (most recent call last)
t)
<ipython-input-66-9248b0c5a70f> in <module>
      6
      7 c1 = SecretKeeper(hidden = "supersecret")
----> 8 c1.hidden

AttributeError: 'SecretKeeper' object has no attribute 'hidden'
```

Can we hide it away somewhat?

In [67]:

```
dir(c1)
```

Out[67]:

```
['_SecretKeeper__hidden',  
 '__class__',  
 '__delattr__',  
 '__dict__',  
 '__dir__',  
 '__doc__',  
 '__eq__',  
 '__format__',  
 '__ge__',  
 '__getattr__',  
 '__gt__',  
 '__hash__',  
 '__init__',  
 '__init_subclass__',  
 '__le__',  
 '__lt__',  
 '__module__',  
 '__ne__',  
 '__new__',  
 '__reduce__',  
 '__reduce_ex__',  
 '__repr__',  
 '__setattr__',  
 '__sizeof__',  
 '__str__',  
 '__subclasshook__',  
 '__weakref__']
```

Can we find it anyway?

In [ ]:

```
# Yes, if you know name mangling. But take it as a signal from the programmer that you  
  shouldn't.
```

```
# Left as exercise.
```

In [ ]:

```
# And yes, you can reach the instance variables even if you don't know name mangling.  
  # Still something to avoid.
```

Conclusion: follow the guidelines. **Assume that non-hidden attributes are public.**

- ...but we can tailor access by `__getattr__` and `__setattr__`.



In [69]:

```
# Bonus task (outside the scope of this course): using custom methods to steer access.

class Mirror:
    def __getattr__(self, name):

        # This function might do something special, eg pass the message on
        # over the internet, to an object on a remote server.

        return "You wanted {}, you say?".format(name)

val = Mirror()
print(val.asdasdasdfasdfasdf)
dir(val) # not a lot here
```

You wanted asdasdasdfasdfasdf, you say?

Out[69]:

[]

## Most decisions are made by the objects

In [74]:

```
# Somewhat contrived example!

class Snake():
    def __add__(self, other):
        return "Snake!"

class Ladder():
    def __add__(self, other):
        return "Hello, this is Ladder."

# Apart from possible issues with inheriting from int, will it commute?
p1 = Snake()
p2 = Ladder()
p2 + p1
```

Out[74]:

'Hello, this is Ladder.'

Why useful to know?

- Python methods carried by objects do the heavy lifting, even when there is no `.method()` in the call.
- Following conventions is good. If you implement your own classes, you might want `p1 + p2 == p2 + p1` to hold, even if Python doesn't force you.
- Builtins are hard to avoid. Shows why your nice `my_clever_vector * 5` might be different from `5 * my_clever_vector`.

## Getters? Setters? Interesting feature: properties

- Controlling getting, assignment, deletion of value.
- What happens if we write `myobj.x = 5` ? Can be controlled via functions, like in eg C#.

In [82]:

```
class GetterTestClass():
    def __init__(self, n = 0):
        self.__n = n

    def __get_n(self):
        return self.__n

    def __set_n(self, new_val):
        if new_val >= 0:
            self.__n = new_val
        else:
            raise ValueError("n must be non-negative!")

    n = property(fget=__get_n, fset = __set_n)

# Left out: fdel. When someone deletes an attribute.

# When someone tries to get n, __get_n will be called. (Etc)

c1 = GetterTestClass(n = 100)
c1.n
c1.n = 3
#c1.n
```

Note: `@property` -syntax also available. You'll do this in lab 3A.

- Slightly weird. Sidesteps the usual `=` always meaning changing labels for class attributes.
- **Know that it's there, don't (ab)use.**

## Useful pointer: dataclasses

- We might structure data by "dummy" classes.
- Predictably: initialisers with the relevant attributes, etc.
  - Prime example: the `Car` above.
- Since Python 3.7, we have [dataclasses](https://docs.python.org/3/library/dataclasses.html) (<https://docs.python.org/3/library/dataclasses.html>) in the standard library. (Also pip-installed in the currently-Python 3.6.3 lab environment if you want to try it.)

## Concepts to explore on your own

- Promises that a class should implement some behaviour (eg "this class should support sequence methods"), that can be checked by the system.

[Abstract Base Classes \(ABC:s\)](https://docs.python.org/3/glossary.html#term-abstract-base-class) (<https://docs.python.org/3/glossary.html#term-abstract-base-class>).

In [ ]:

```
import collections
isinstance([1,2,3], collections.abc.Sequence)
```

See also the `abc` module in standard library.

- Mixins. Adding a capability to your class (note: here we include implementation). See Lutz.

In [83]:

```
# We can stick the capability to return five (or something more useful) into  
# a class by just making it a subclass of ReturnFiveMixin.
```

```
class ReturnFiveMixin:
    def return_five(self):
        return 5

class MyClass(ReturnFiveMixin):
    # No code here!
    pass
```

```
obj = MyClass()
obj.return_five()
```

Out[83]:

5

- Class and instance attribute conventions.

In [86]:

```
class MyVal:
    common_to_all_myvals = 99
    def __init__(self, val):
        self.val = val

one_val = MyVal(val = 1)
two_val = MyVal(val = 2)
one_val.val # in the instance
one_val.common_to_all_myvals # in the class
```

Out[86]:

99

In [88]:

```
one_val.common_to_all_myvals = "new value. Is this only in one_val?"
two_val.common_to_all_myvals # Actually, just a local name in one_val.
```

Out[88]:

99

Note: you may have class methods which are reached via `<Class>.method()` , as we might reach `MyVal.common_to_all_myvals` above.

## Other interesting *language* details

- Decorators. Transforming Python functions. (Telltale sign in code: the `@` sign. `@something` , such as `@property` , `@dataclass` .)
- Annotations. [PEP3107](https://www.python.org/dev/peps/pep-3107/) (<https://www.python.org/dev/peps/pep-3107/>)

In [89]:

```
def sq(n : int,
      otherarg : "annotations can be anything" = ""):
    """Return the value n^2."""
    return n*n

sq.__annotations__    # can be used by some code analysis tools, documentation tools et
c.
```

Out[89]:

```
{'n': int, 'otherarg': 'annotations can be anything'}
```

- Consequence: type information might be useful eg
  - for code analysis, See eg [mypy](http://www.mypy-lang.org/) (<http://www.mypy-lang.org/>).
  - IDE:s.

## Sets (and built-in types)

- Why do we care?
- So far:
  - `list` . Mutable, quick access by index. Finding in general slow.
  - `tuple` . Immutable. Fast, efficient, quick access by index. But: finding in general slow.
  - `dict` . Mapping, based on hash tables. Very fast access by key, finding by value slow. Membership test very fast. Slight memory overhead. Requires values to be hashable.
    - Errata: insertion-order iteration since Python 3.7 (not *unordered*).
- Yet another useful type: `set` . Based on hash tables, with **very fast membership tests**. Mathematical-set methods.

In [96]:

```
import profile, random

N = 9999999
vals_list = list(range(N))
vals_set = set(range(N))

# Pick some random element, just for demonstration purposes.
needle = random.randint(0, N)

print("--- vals_list")
profile.run(str(needle) + " in vals_list")

print("--- vals_set")
profile.run(str(needle) + " in vals_set")
```

```
--- vals_list
    4 function calls in 0.079 seconds
```

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.079	0.079	:0(exec)
1	0.000	0.000	0.000	0.000	:0(setprofile)
1	0.079	0.079	0.079	0.079	<string>:1(<module>)
1	0.000	0.000	0.079	0.079	profile:0(7884751 in vals_li
st)	0	0.000	0.000		profile:0(profiler)

```
--- vals_set
    4 function calls in 0.000 seconds
```

Ordered by: standard name

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.000	0.000	:0(exec)
1	0.000	0.000	0.000	0.000	:0(setprofile)
1	0.000	0.000	0.000	0.000	<string>:1(<module>)
1	0.000	0.000	0.000	0.000	profile:0(7884751 in vals_se
t)	0	0.000	0.000		profile:0(profiler)

Conclusion: if you need to perform lots of lookups based on the keys, sets, dicts etc might be useful.

In [98]:

```
words = set(["cat", "snape", "doge"])
animals = set(["cat"])

# What are the commonalities?
words.intersection(animals)
```

Out[98]:

{'cat'}

In [99]:

```
# Which words are not in animals?
words.difference(animals)
```

Out[99]:

```
{'doge', 'snape'}
```

In [103]:

```
# How would we add a word?
animals.add("giraffe")
animals.add("giraffe") # Redundant - sets contain no duplicates.
animals
```

Out[103]:

```
{'cat', 'giraffe'}
```

- Immutable version: frozenset . Can be used as a key in a dictionary.

In [ ]:

## Aside: modules

- Those `import math`, `import my_own_module` ...
- Where do they come from? [[reference \(https://docs.python.org/3.7/tutorial/modules.html#the-module-search-path\)](https://docs.python.org/3.7/tutorial/modules.html#the-module-search-path)]
  - Beginner's note: some installation system usually takes care of this for you. But useful to know if something breaks.
- When you write `my_module.py` you can do `import my_module` at least from the same directory (see search path above).
- You will see `__init__.py` around. This concerns [packages \(https://docs.python.org/3.7/tutorial/modules.html#packages\)](https://docs.python.org/3.7/tutorial/modules.html#packages).
- Namespaces.

In [93]:

```
# Example package: sklearn, and its datasets  
import sklearn  
import sklearn.datasets  
help(sklearn.datasets)
```

Help on package sklearn.datasets in sklearn:

## NAME

sklearn.datasets

## DESCRIPTION

The :mod:`sklearn.datasets` module includes utilities to load datasets, including methods to load and fetch popular reference datasets. It also features some artificial data generators.

## PACKAGE CONTENTS

\_svmlight\_format  
base  
california\_housing  
covtype  
kddcup99  
lfw  
mlcomp  
mldata  
olivetti\_faces  
openml  
rcv1  
samples\_generator  
setup  
species\_distributions  
svmlight\_format  
tests (package)  
twenty\_newsgroups

## FUNCTIONS

clear\_data\_home(data\_home=None)  
Delete all the content of the data home cache.

Parameters

-----

data\_home : str | None  
The path to scikit-learn data dir.

dump\_svmlight\_file(X, y, f, zero\_based=True, comment=None, query\_id=None, multilabel=False)

Dump the dataset in svmlight / libsvm file format.

This format is a text-based format, with one sample per line. It does not store zero valued features hence is suitable for sparse datasets.

The first element of each line can be used to store a target variable to predict.

Parameters

-----

X : {array-like, sparse matrix}, shape = [n\_samples, n\_features]  
Training vectors, where n\_samples is the number of samples and n\_features is the number of features.

y : {array-like, sparse matrix}, shape = [n\_samples, n\_labels]  
Target values. Class labels must be an



integer or float, or array-like objects of integer or float for  
 multilabel classifications.

**f** : string or file-like in binary mode  
 If string, specifies the path that will contain the data.  
 If file-like, data will be written to f. f should be opened in  
 binary mode.

**zero\_based** : boolean, optional  
 Whether column indices should be written zero-based (True) or  
 one-based (False).

**comment** : string, optional  
 Comment to insert at the top of the file. This should be either  
 a Unicode string, which will be encoded as UTF-8, or an ASCII byte  
 string.  
 If a comment is given, then it will be preceded by one that identifies  
 the file as having been dumped by scikit-learn. Note that not all  
 tools grok comments in SVMlight files.

**query\_id** : array-like, shape = [n\_samples]  
 Array containing pairwise preference constraints (qid in svmli  
 ght format).

**multilabel** : boolean, optional  
 Samples may have several labels each (see  
<https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/multilabel.html>)

.. versionadded:: 0.17  
 parameter \*multilabel\* to support multilabel datasets.

**fetch\_20newsgroups**(data\_home=None, subset='train', categories=None, shuffle=True, random\_state=42, remove=(), download\_if\_missing=True)  
 Load the filenames and data from the 20 newsgroups dataset (classification).

Download it if necessary.

```

=====
Classes                20
Samples total          18846
Dimensionality          1
Features                text
=====
```

Read more in the :ref:`User Guide <20newsgroups\_dataset>`.

Parameters

-----

**data\_home** : optional, default: None  
 Specify a download and cache folder for the datasets. If None, all scikit-learn data is stored in '~/scikit\_learn\_data' subfo

lders.

subset : 'train' or 'test', 'all', optional  
 Select the dataset to load: 'train' for the training set, 'test' for the test set, 'all' for both, with shuffled ordering.

categories : None or collection of string or unicode  
 If None (default), load all the categories.  
 If not None, list of category names to load (other categories ignored).

shuffle : bool, optional  
 Whether or not to shuffle the data: might be important for models that make the assumption that the samples are independent and identically distributed (i.i.d.), such as stochastic gradient descent.

random\_state : int, RandomState instance or None (default)  
 Determines random number generation for dataset shuffling. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

remove : tuple  
 May contain any subset of ('headers', 'footers', 'quotes'). Each of these are kinds of text that will be detected and removed from the newsgroup posts, preventing classifiers from overfitting on metadata.  
 'headers' removes newsgroup headers, 'footers' removes blocks at the ends of posts that look like signatures, and 'quotes' removes lines that appear to be quoting another post.  
 'headers' follows an exact standard; the other filters are not always correct.

download\_if\_missing : optional, True by default  
 If False, raise an IOError if the data is not locally available instead of trying to download the data from the source site.

Returns

-----

bunch : Bunch object  
 bunch.data: list, length [n\_samples]  
 bunch.target: array, shape [n\_samples]  
 bunch filenames: list, length [n\_classes]  
 bunch.DESCR: a description of the dataset.

fetch\_20newsgroups\_vectorized(subset='train', remove=(), data\_home=None, download\_if\_missing=True, return\_X\_y=False)  
 Load the 20 newsgroups dataset and vectorize it into token counts (classification).

Download it if necessary.

This is a convenience function; the transformation is done using t

he

default settings for

:class:`sklearn.feature\_extraction.text.CountVectorizer`. For more advanced usage (stopword filtering, n-gram extraction, etc.), comb

ine

fetch\_20newsgroups with a custom

```
:class:`sklearn.feature_extraction.text.CountVectorizer`,
:class:`sklearn.feature_extraction.text.HashingVectorizer`,
:class:`sklearn.feature_extraction.text.TfidfTransformer` or
:class:`sklearn.feature_extraction.text.TfidfVectorizer`.
```

```
=====
Classes                20
Samples total          18846
Dimensionality         130107
Features               real
=====
```

Read more in the :ref:`User Guide <20newsgroups\_dataset>`.

Parameters

-----

subset : 'train' or 'test', 'all', optional

Select the dataset to load: 'train' for the training set, 'tes

t'

for the test set, 'all' for both, with shuffled ordering.

remove : tuple

May contain any subset of ('headers', 'footers', 'quotes'). Ea

ch of

these are kinds of text that will be detected and removed from

the

newsgroup posts, preventing classifiers from overfitting on metadata.

'headers' removes newsgroup headers, 'footers' removes blocks

at the

ends of posts that look like signatures, and 'quotes' removes

lines

that appear to be quoting another post.

data\_home : optional, default: None

Specify an download and cache folder for the datasets. If Non

e,

all scikit-learn data is stored in '~/scikit\_learn\_data' subfo

lders.

download\_if\_missing : optional, True by default

If False, raise an IOError if the data is not locally availabl

e

instead of trying to download the data from the source site.

return\_X\_y : boolean, default=False.

If True, returns ``(data.data, data.target)`` instead of a Bun

ch

object.

.. versionadded:: 0.20

## Returns

-----

bunch : Bunch object

bunch.data: sparse matrix, shape [n\_samples, n\_features]

bunch.target: array, shape [n\_samples]

bunch.target\_names: list, length [n\_classes]

bunch.DESCR: a description of the dataset.

(data, target) : tuple if ``return\_X\_y`` is True

.. versionadded:: 0.20

fetch\_california\_housing(data\_home=None, download\_if\_missing=True, return\_X\_y=False)

Load the California housing dataset (regression).

```

=====
Samples total      20640
Dimensionality      8
Features           real
Target             real 0.15 - 5.
=====

```

Read more in the :ref:`User Guide &lt;california\_housing\_dataset&gt;`.

## Parameters

-----

data\_home : optional, default: None

Specify another download and cache folder for the datasets. By default all scikit-learn data is stored in '~/.scikit\_learn\_data' subfolders.

download\_if\_missing : optional, default=True

If False, raise a IOError if the data is not locally available instead of trying to download the data from the source site.

return\_X\_y : boolean, default=False.

If True, returns ``(data.data, data.target)`` instead of a Bunch object.

.. versionadded:: 0.20

## Returns

-----

dataset : dict-like object with the following attributes:

dataset.data : ndarray, shape [20640, 8]

Each row corresponding to the 8 feature values in order.

dataset.target : numpy array of shape (20640,)

Each value corresponds to the average house value in units of 100,000.

dataset.feature\_names : array of length 8

Array of ordered feature names used in the dataset.

dataset.DESCR : string

## Description of the California housing dataset.

```
(data, target) : tuple if ``return_X_y`` is True
```

```
.. versionadded:: 0.20
```

## Notes

```
-----
```

This dataset consists of 20,640 samples and 9 features.

```
fetch_covtype(data_home=None, download_if_missing=True, random_state=None,
               shuffle=False, return_X_y=False)
```

Load the covtype dataset (classification).

Download it if necessary.

```
=====
Classes                7
Samples total          581012
Dimensionality          54
Features                int
=====
```

Read more in the :ref:`User Guide <covtype\_dataset>`.

## Parameters

```
-----
```

`data_home` : string, optional  
Specify another download and cache folder for the datasets. By default all scikit-learn data is stored in '~/.scikit\_learn\_data' subfolders.

`download_if_missing` : boolean, default=True  
If False, raise a IOError if the data is not locally available instead of trying to download the data from the source site.

`random_state` : int, RandomState instance or None (default)  
Determines random number generation for dataset shuffling. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

`shuffle` : bool, default=False  
Whether to shuffle dataset.

`return_X_y` : boolean, default=False.  
If True, returns ``(data.data, data.target)`` instead of a Bunch object.

```
.. versionadded:: 0.20
```

## Returns

```
-----
```

`dataset` : dict-like object with the following attributes:

`dataset.data` : numpy array of shape (581012, 54)  
Each row corresponds to the 54 features in the dataset.

```

dataset.target : numpy array of shape (581012,)
    Each value corresponds to one of the 7 forest covertypes with
values
    ranging between 1 to 7.

dataset.DESCR : string
    Description of the forest covertype dataset.

(data, target) : tuple if ``return_X_y`` is True

.. versionadded:: 0.20

fetch_kddcup99(subset=None, data_home=None, shuffle=False, random_state=None, percent10=True, download_if_missing=True, return_X_y=False)
    Load the kddcup99 dataset (classification).

    Download it if necessary.

=====
Classes                23
Samples total          4898431
Dimensionality         41
Features               discrete (int) or continuous (float)
=====

Read more in the :ref:`User Guide <kddcup99_dataset>`.

.. versionadded:: 0.18

Parameters
-----
subset : None, 'SA', 'SF', 'http', 'smtp'
    To return the corresponding classical subsets of kddcup 99.
    If None, return the entire kddcup 99 dataset.

data_home : string, optional
    Specify another download and cache folder for the datasets. By
default
    all scikit-learn data is stored in '~/scikit_learn_data' subfo
lders.

.. versionadded:: 0.19

shuffle : bool, default=False
    Whether to shuffle dataset.

random_state : int, RandomState instance or None (default)
    Determines random number generation for dataset shuffling and
for
    selection of abnormal samples if `subset='SA'`. Pass an int fo
r
    reproducible output across multiple function calls.
    See :term:`Glossary <random_state>`.

percent10 : bool, default=True
    Whether to load only 10 percent of the data.

download_if_missing : bool, default=True
    If False, raise a IOError if the data is not locally available
    instead of trying to download the data from the source site.

return_X_y : boolean, default=False.

```

If True, returns ``(data, target)`` instead of a Bunch object.

See

below for more information about the `data` and `target` objec

t.

.. versionadded:: 0.20

Returns

-----

data : Bunch

Dictionary-like object, the interesting attributes are:

- 'data', the data to learn.
- 'target', the regression target for each sample.
- 'DESCR', a description of the dataset.

(data, target) : tuple if ``return\_X\_y`` is True

.. versionadded:: 0.20

fetch\_lfw\_pairs(subset='train', data\_home=None, funneled=True, resize=0.5, color=False, slice\_=(slice(70, 195, None), slice(78, 172, None)), download\_if\_missing=True)

Load the Labeled Faces in the Wild (LFW) pairs dataset (classification).

Download it if necessary.

```
=====
Classes                               5749
Samples total                         13233
Dimensionality                        5828
Features                             real, between 0 and 255
=====
```

In the official `README.txt`\_ this task is described as the "Restricted" task. As I am not sure as to implement the "Unrestricted" variant correctly, I left it as unsupported for now.

w.

.. \_`README.txt`: <http://vis-www.cs.umass.edu/lfw/README.txt>

The original images are 250 x 250 pixels, but the default slice and resize

arguments reduce them to 62 x 47.

Read more in the :ref:`User Guide <labeled\_faces\_in\_the\_wild\_dataset>`.

Parameters

-----

subset : optional, default: 'train'

Select the dataset to load: 'train' for the development training

set, 'test' for the development test set, and '10\_folds' for the

official evaluation set that is meant to be used with a 10-fold

cross validation.

data\_home : optional, default: None

Specify another download and cache folder for the datasets. By

default all scikit-learn data is stored in '~/scikit\_learn\_data'

subfolders.

funneled : boolean, optional, default: True  
Download and use the funneled variant of the dataset.

resize : float, optional, default 0.5  
Ratio used to resize the each face picture.

color : boolean, optional, default False  
Keep the 3 RGB channels instead of averaging them to a single gray level channel. If color is True the shape of the data has one more dimension than the shape with color = False.

slice\_ : optional  
Provide a custom 2D slice (height, width) to extract the 'interesting' part of the jpeg files and avoid use statistical correlation from the background

download\_if\_missing : optional, True by default  
If False, raise a IOError if the data is not locally available instead of trying to download the data from the source site.

#### Returns

-----

The data is returned as a Bunch object with the following attributes:

data : numpy array of shape (2200, 5828). Shape depends on ``subset``.

Each row corresponds to 2 ravel'd face images of original size 62 x 47 pixels. Changing the ``slice\_``, ``resize`` or ``subset`` parameters will change the shape of the output.

pairs : numpy array of shape (2200, 2, 62, 47). Shape depends on ``subset``

Each row has 2 face images corresponding to same or different person from the dataset containing 5749 people. Changing the ``slice\_``, ``resize`` or ``subset`` parameters will change the shape of the output.

target : numpy array of shape (2200,). Shape depends on ``subset``.

Labels associated to each pair of images. The two label values being different persons or the same person.

DESCR : string  
Description of the Labeled Faces in the Wild (LFW) dataset.

fetch\_lfw\_people(data\_home=None, funneled=True, resize=0.5, min\_faces\_per\_person=0, color=False, slice\_=(slice(70, 195, None), slice(78, 172, None)), download\_if\_missing=True, return\_X\_y=False)  
Load the Labeled Faces in the Wild (LFW) people dataset (classification).



Download it if necessary.

```
=====
Classes                      5749
Samples total                13233
Dimensionality               5828
Features                    real, between 0 and 255
=====
```

Read more in the :ref:`User Guide <labeled\_faces\_in\_the\_wild\_data  
et>`.

#### Parameters

-----

`data_home` : optional, default: None  
Specify another download and cache folder for the datasets. By default all scikit-learn data is stored in '~/.scikit\_learn\_data' subfolders.

`funneled` : boolean, optional, default: True  
Download and use the funneled variant of the dataset.

`resize` : float, optional, default 0.5  
Ratio used to resize the each face picture.

`min_faces_per_person` : int, optional, default None  
The extracted dataset will only retain pictures of people that have at least `min_faces_per_person` different pictures.

`color` : boolean, optional, default False  
Keep the 3 RGB channels instead of averaging them to a single gray level channel. If color is True the shape of the data has one more dimension than the shape with color = False.

`slice_` : optional  
Provide a custom 2D slice (height, width) to extract the 'interesting' part of the jpeg files and avoid use statistical correlation from the background

`download_if_missing` : optional, True by default  
If False, raise a IOError if the data is not locally available instead of trying to download the data from the source site.

`return_X_y` : boolean, default=False.  
If True, returns ``(dataset.data, dataset.target)`` instead of a Bunch object. See below for more information about the `dataset.data` and `dataset.target` object.

.. versionadded:: 0.20

#### Returns

-----

`dataset` : dict-like object with the following attributes:

`dataset.data` : numpy array of shape (13233, 2914)  
Each row corresponds to a ravelled face image of original size

62 x 47

pixels. Changing the ``slice`` or resize parameters will change the shape of the output.

dataset.images : numpy array of shape (13233, 62, 47)  
Each row is a face image corresponding to one of the 5749 people in the dataset. Changing the ``slice`` or resize parameters will change the shape of the output.

dataset.target : numpy array of shape (13233,)  
Labels associated to each face image. Those labels range from 0-5748 and correspond to the person IDs.

dataset.DESCR : string  
Description of the Labeled Faces in the Wild (LFW) dataset.

(data, target) : tuple if ``return\_X\_y`` is True

.. versionadded:: 0.20

fetch\_mldata(dataname, target\_name='label', data\_name='data', transpose\_data=True, data\_home=None)  
DEPRECATED: fetch\_mldata was deprecated in version 0.20 and will be removed in version 0.22

Fetch an mldata.org data set

mldata.org is no longer operational.

If the file does not exist yet, it is downloaded from mldata.org .

mldata.org does not have an enforced convention for storing data or naming the columns in a data set. The default behavior of this function works well with the most common cases:

- 1) data values are stored in the column 'data', and target values in the column 'label'
- 2) alternatively, the first column stores target values, and the second data values
- 3) the data array is stored as ``n\_features x n\_samples`` , and thus needs to be transposed to match the ``sklearn`` standard

Keyword arguments allow to adapt these defaults to specific datasets (see parameters ``target\_name``, ``data\_name``, ``transpose\_data``, and the examples below).

mldata.org data sets may have multiple columns, which are stored in the Bunch object with their original name.

.. deprecated:: 0.20  
Will be removed in version 0.22

#### Parameters

-----

**dataname** : str  
Name of the data set on mldata.org,  
e.g.: "leukemia", "Whistler Daily Snowfall", etc.  
The raw name is automatically converted to a mldata.org UR

L .

**target\_name** : optional, default: 'label'  
Name or index of the column containing the target values.

**data\_name** : optional, default: 'data'  
Name or index of the column containing the data.

**transpose\_data** : optional, default: True  
If True, transpose the downloaded data array.

**data\_home** : optional, default: None  
Specify another download and cache folder for the data set

s. By default

all scikit-learn data is stored in '~/scikit\_learn\_data' subfolders.

#### Returns

-----

**data** : Bunch  
Dictionary-like object, the interesting attributes are:  
'data', the data to learn, 'target', the classification labels,  
'DESCR', the full description of the dataset, and  
'COL\_NAMES', the original names of the dataset columns.

**fetch\_olivetti\_faces**(data\_home=None, shuffle=False, random\_state=0, download\_if\_missing=True)  
Load the Olivetti faces data-set from AT&T (classification).

Download it if necessary.

=====	=====
Classes	40
Samples total	400
Dimensionality	4096
Features	real, between 0 and 1
=====	=====

Read more in the :ref:`User Guide <olivetti\_faces\_dataset>`.

#### Parameters

-----

**data\_home** : optional, default: None  
Specify another download and cache folder for the datasets. By default  
all scikit-learn data is stored in '~/scikit\_learn\_data' subfolders.

shuffle : boolean, optional

If True the order of the dataset is shuffled to avoid having images of the same person grouped.

random\_state : int, RandomState instance or None (default=0)

Determines random number generation for dataset shuffling. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

download\_if\_missing : optional, True by default

If False, raise a IOError if the data is not locally available instead of trying to download the data from the source site.

Returns

-----

An object with the following attributes:

data : numpy array of shape (400, 4096)

Each row corresponds to a ravelled face image of original size 64 x 64 pixels.

images : numpy array of shape (400, 64, 64)

Each row is a face image corresponding to one of the 40 subjects of the dataset.

target : numpy array of shape (400, )

Labels associated to each face image. Those labels are ranging from 0-39 and correspond to the Subject IDs.

DESCR : string

Description of the modified Olivetti Faces Dataset.

fetch\_openml(name=None, version='active', data\_id=None, data\_home=None, target\_column='default-target', cache=True, return\_X\_y=False)

Fetch dataset from openml by name or dataset id.

Datasets are uniquely identified by either an integer ID or by a combination of name and version (i.e. there might be multiple versions of the 'iris' dataset). Please give either name or data\_id

(not both). In case a name is given, a version can also be provided.

Read more in the :ref:`User Guide <openml>`.

.. note:: EXPERIMENTAL

The API is experimental in version 0.20 (particularly the return value structure), and might have small backward-incompatible changes in future releases.

Parameters

-----

name : str or None

String identifier of the dataset. Note that OpenML can have multiple

datasets with the same name.

`version : integer or 'active', default='active'`

Version of the dataset. Can only be provided if also `name`

is given.

If 'active' the oldest version that's still active is used. Si

nce

there may be more than one active version of a dataset, and th

ose

versions may fundamentally be different from one another, sett

ing an

exact version is highly recommended.

`data_id : int or None`

OpenML ID of the dataset. The most specific way of retrieving

a

dataset. If `data_id` is not given, `name` (and potential version)

are

used to obtain a dataset.

`data_home : string or None, default None`

Specify another download and cache folder for the data sets. B

y default

all scikit-learn data is stored in `~/scikit_learn_data` subfo

lders.

`target_column : string, list or None, default 'default-target'`

Specify the column name in the data to use as target. If

'default-target', the standard target column a stored on the s

erver

is used. If `None`, all columns are returned as data and the  
target is `None`. If list (of strings), all columns with the

se names

are returned as multi-target (Note: not all scikit-learn class

ifiers

can handle all types of multi-output combinations)

`cache : boolean, default=True`

Whether to cache downloaded datasets using joblib.

`return_X_y : boolean, default=False.`

If True, returns `(data, target)` instead of a Bunch object.

See

below for more information about the `data` and `target` objec

ts.

Returns

-----

`data : Bunch`

Dictionary-like object, with attributes:

`data : np.array or scipy.sparse.csr_matrix of floats`

The feature matrix. Categorical features are encoded as or

dinals.

`target : np.array`

The regression target or classification labels, if applica

ble.

Dtype is float if numeric, and object if categorical.

`DESCR : str`

The full description of the dataset

```

feature_names : list
    The names of the dataset columns
categories : dict
    Maps each categorical feature name to a list of values, such
    that the value encoded as i is ith in the list.
details : dict
    More metadata from OpenML

```

```
(data, target) : tuple if ``return_X_y`` is True
```

```
.. note:: EXPERIMENTAL
```

This interface is **experimental** as at version 0.20 and subsequent releases may change attributes without notice (although there should only be minor changes to ``data`` and ``target``).

Missing values in the 'data' are represented as NaN's. Missing values in 'target' are represented as NaN's (numerical target) or None (categorical target)

```

fetch_rcv1(data_home=None, subset='all', download_if_missing=True, random_state=None, shuffle=False, return_X_y=False)
    Load the RCV1 multilabel dataset (classification).

```

Download it if necessary.

Version: RCV1-v2, vectors, full sets, topics multilabels.

```

=====
Classes                                103
Samples total                          804414
Dimensionality                         47236
Features                               real, between 0 and 1
=====

```

Read more in the :ref:`User Guide <rcv1\_dataset>`.

```
.. versionadded:: 0.17
```

Parameters

-----

**data\_home** : string, optional  
Specify another download and cache folder for the datasets. By default all scikit-learn data is stored in '~/.scikit\_learn\_data' subfolders.

**subset** : string, 'train', 'test', or 'all', default='all'  
Select the dataset to load: 'train' for the training set (23149 samples), 'test' for the test set (781265 samples), 'all' for both, with the training samples first if shuffle is False.  
This follows the official LYRL2004 chronological split.

**download\_if\_missing** : boolean, default=True  
If False, raise a IOError if the data is not locally available instead of trying to download the data from the source site.

random\_state : int, RandomState instance or None (default)  
 Determines random number generation for dataset shuffling. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

shuffle : bool, default=False  
 Whether to shuffle dataset.

return\_X\_y : boolean, default=False.  
 If True, returns ``(dataset.data, dataset.target)`` instead of a Bunch object. See below for more information about the `dataset.data` and `dataset.target` object.

.. versionadded:: 0.20

Returns

-----

dataset : dict-like object with the following attributes:

dataset.data : scipy csr array, dtype np.float64, shape (804414, 4)  
 The array has 0.16% of non zero values.

dataset.target : scipy csr array, dtype np.uint8, shape (804414, 1)  
 Each sample has a value of 1 in its categories, and 0 in other categories. The array has 3.15% of non zero values.

dataset.sample\_id : numpy array, dtype np.uint32, shape (804414,)  
 Identification number of each sample, as ordered in dataset.data.

dataset.target\_names : numpy array, dtype object, length (103)  
 Names of each target (RCV1 topics), as ordered in dataset.target.

dataset.DESCR : string  
 Description of the RCV1 dataset.

(data, target) : tuple if ``return\_X\_y`` is True

.. versionadded:: 0.20

fetch\_species\_distributions(data\_home=None, download\_if\_missing=True)  
 Loader for species distribution dataset from Phillips et. al. (2006)

Read more in the :ref:`User Guide <datasets>`.

Parameters

-----

data\_home : optional, default: None  
 Specify another download and cache folder for the datasets. By default all scikit-learn data is stored in '~/.scikit\_learn\_data' subfolders.

download\_if\_missing : optional, True by default  
 If False, raise a IOError if the data is not locally available  
 instead of trying to download the data from the source site.

Returns

-----

The data is returned as a Bunch object with the following attributes:

coverages : array, shape = [14, 1592, 1212]

These represent the 14 features measured at each point of the map grid.

The latitude/longitude values for the grid are discussed below.

Missing data is represented by the value -9999.

train : record array, shape = (1624,)

The training points for the data. Each point has three fields:

- train['species'] is the species name
- train['dd long'] is the longitude, in degrees
- train['dd lat'] is the latitude, in degrees

test : record array, shape = (620,)

The test points for the data. Same format as the training data.

Nx, Ny : integers

The number of longitudes (x) and latitudes (y) in the grid

x\_left\_lower\_corner, y\_left\_lower\_corner : floats

The (x,y) position of the lower-left corner, in degrees

grid\_size : float

The spacing between points of the grid, in degrees

References

-----

\* "Maximum entropy modeling of species geographic distributions"

<<http://rob.schapire.net/papers/ecolmod.pdf>>\_

S. J. Phillips, R. P. Anderson, R. E. Schapire - Ecological Modelling, 190:231-259, 2006.

Notes

-----

This dataset represents the geographic distribution of species.

The dataset is provided by Phillips et. al. (2006).

The two species are:

- "Bradypus variegatus"

<<http://www.iucnredlist.org/details/3038/0>>\_ ,  
 the Brown-throated Sloth.

- "Microryzomys minutus"

<<http://www.iucnredlist.org/details/13408/0>>\_ ,



also known as the Forest Small Rice Rat, a rodent that lives in  
Peru,  
Colombia, Ecuador, Peru, and Venezuela.

- For an example of using this dataset with scikit-learn, see  
:ref:`examples/applications/plot\_species\_distribution\_modeling.py`

```
<sphinx_gallery_auto_examples_applications_plot_species_distribution_modeling.py>`.
```

```
get_data_home(data_home=None)
    Return the path of the scikit-learn data dir.
```

This folder is used by some large dataset loaders to avoid downloading the  
data several times.

By default the data dir is set to a folder named 'scikit\_learn\_data' in the  
user home folder.

Alternatively, it can be set by the 'SCIKIT\_LEARN\_DATA' environment  
variable or programmatically by giving an explicit folder path. The  
symbol '~' is expanded to the user home folder.

If the folder does not already exist, it is automatically created.

Parameters

-----

data\_home : str | None  
 The path to scikit-learn data dir.

```
load_boston(return_X_y=False)
    Load and return the boston house-prices dataset (regression).
```

```
=====
Samples total      506
Dimensionality     13
Features           real, positive
Targets            real 5. - 50.
=====
```

Read more in the :ref:`User Guide <boston\_dataset>`.

Parameters

-----

return\_X\_y : boolean, default=False.  
 If True, returns ``(data, target)`` instead of a Bunch object.  
 See below for more information about the `data` and `target` objects.

object.

.. versionadded:: 0.18

Returns

-----

data : Bunch  
 Dictionary-like object, the interesting attributes are:  
 'data', the data to learn, 'target', the regression targets,  
 'DESCR', the full description of the dataset,

and 'filename', the physical location of boston csv dataset (added in version `0.20`).

```
(data, target) : tuple if ``return_X_y`` is True
```

```
.. versionadded:: 0.18
```

#### Notes

```
-----
```

```
.. versionchanged:: 0.20
    Fixed a wrong data point at [445, 0].
```

#### Examples

```
-----
```

```
>>> from sklearn.datasets import load_boston
>>> boston = load_boston()
>>> print(boston.data.shape)
(506, 13)
```

```
load_breast_cancer(return_X_y=False)
```

Load and return the breast cancer wisconsin dataset (classification).

The breast cancer dataset is a classic and very easy binary classification dataset.

```
=====
Classes                2
Samples per class      212(M),357(B)
Samples total          569
Dimensionality         30
Features               real, positive
=====
```

Read more in the :ref:`User Guide <breast\_cancer\_dataset>`.

#### Parameters

```
-----
```

```
return_X_y : boolean, default=False
```

If True, returns ``(data, target)`` instead of a Bunch object. See below for more information about the `data` and `target` o

bject.

```
.. versionadded:: 0.18
```

#### Returns

```
-----
```

```
data : Bunch
```

Dictionary-like object, the interesting attributes are: 'data', the data to learn, 'target', the classification label

s,

'target\_names', the meaning of the labels, 'feature\_names', the

e

meaning of the features, and 'DESCR', the full description of the dataset, 'filename', the physical location of breast cancer csv dataset (added in version `0.20`).

```
(data, target) : tuple if ``return_X_y`` is True
```

```
.. versionadded:: 0.18
```

The copy of UCI ML Breast Cancer Wisconsin (Diagnostic) dataset is downloaded from:

<https://goo.gl/U2Uwz2>

#### Examples

-----

Let's say you are interested in the samples 10, 50, and 85, and want to know their class name.

```
>>> from sklearn.datasets import load_breast_cancer
>>> data = load_breast_cancer()
>>> data.target[[10, 50, 85]]
array([0, 1, 0])
>>> list(data.target_names)
['malignant', 'benign']
```

`load_diabetes(return_X_y=False)`

Load and return the diabetes dataset (regression).

=====	=====
Samples total	442
Dimensionality	10
Features	real, $-0.2 < x < 0.2$
Targets	integer 25 - 346
=====	=====

Read more in the :ref:`User Guide <diabetes\_dataset>`.

#### Parameters

-----

`return_X_y` : boolean, default=False.

If True, returns ``(data, target)`` instead of a Bunch object. See below for more information about the ``data`` and ``target`` objects.

object.

.. versionadded:: 0.18

#### Returns

-----

`data` : Bunch

Dictionary-like object, the interesting attributes are:

each sample, `'data_filename'`, the physical location of diabetes data csv dataset, and `'target_filename'`, the physical location of diabetes targets csv dataset (added in version ``0.20``).

`(data, target)` : tuple if `return_X_y` is True

.. versionadded:: 0.18

`load_digits(n_class=10, return_X_y=False)`

Load and return the digits dataset (classification).

Each datapoint is a 8x8 image of a digit.

=====

Classes	10
Samples per class	~180
Samples total	1797
Dimensionality	64
Features	integers 0-16
=====	=====

Read more in the :ref:`User Guide <digits\_dataset>`.

#### Parameters

-----

`n_class` : integer, between 0 and 10, optional (default=10)  
The number of classes to return.

`return_X_y` : boolean, default=False.  
If True, returns ``(data, target)`` instead of a Bunch object.  
See below for more information about the ``data`` and ``target`` o

bject.

.. versionadded:: 0.18

#### Returns

-----

`data` : Bunch  
Dictionary-like object, the interesting attributes are:  
'data', the data to learn, 'images', the images corresponding  
to each sample, 'target', the classification labels for each  
sample, 'target\_names', the meaning of the labels, and 'DESC

R',

the full description of the dataset.

(data, target) : tuple if ``return\_X\_y`` is True

.. versionadded:: 0.18

This is a copy of the test set of the UCI ML hand-written digits datasets

<http://archive.ics.uci.edu/ml/datasets/Optical+Recognition+of+Hand+written+Digits>

#### Examples

-----

To load the data and visualize the images::

```
>>> from sklearn.datasets import load_digits
>>> digits = load_digits()
>>> print(digits.data.shape)
(1797, 64)
>>> import matplotlib.pyplot as plt #doctest: +SKIP
>>> plt.gray() #doctest: +SKIP
>>> plt.matshow(digits.images[0]) #doctest: +SKIP
>>> plt.show() #doctest: +SKIP
```

`load_files(container_path, description=None, categories=None, load_content=True, shuffle=True, encoding=None, decode_error='strict', random_state=0)`

Load text files with categories as subfolder names.

Individual samples are assumed to be files stored a two levels folder

structure such as the following:

```

    container_folder/
        category_1_folder/
            file_1.txt
            file_2.txt
            ...
            file_42.txt
        category_2_folder/
            file_43.txt
            file_44.txt
            ...

```

The folder names are used as supervised signal label names. The individual file names are not important.

This function does not try to extract features into a numpy array or scipy sparse matrix. In addition, if `load_content` is false it does not try to load the files in memory.

To use text files in a scikit-learn classification or clustering algorithm, you will need to use the ``sklearn.feature_extraction.text`` module to build a feature extraction transformer that suits your problem.

If you set `load_content=True`, you should also specify the encoding of the text using the `'encoding'` parameter. For many modern text files, `'utf-8'` will be the correct encoding. If you leave encoding equal to `None`, then the content will be made of bytes instead of Unicode, and you will not be able to use most functions in ``sklearn.feature_extraction.text``.

Similar feature extractors should be built for other kind of unstructured data input such as images, audio, video, ...

Read more in the :ref:`User Guide <datasets>`.

#### Parameters

-----

`container_path` : string or unicode  
Path to the main folder holding one subfolder per category

`description` : string or unicode, optional (default=None)  
A paragraph describing the characteristic of the dataset: its source, reference, etc.

`categories` : A collection of strings or None, optional (default=None)  
If None (default), load all the categories. If not None, list of category names to load (other categories ignored).

`load_content` : boolean, optional (default=True)

Whether to load or not the content of the different files. If true a 'data' attribute containing the text information is present in the data structure returned. If not, a filenames attribute gives the path to the files.

shuffle : bool, optional (default=True)  
Whether or not to shuffle the data: might be important for models that make the assumption that the samples are independent and identically distributed (i.i.d.), such as stochastic gradient descent.

encoding : string or None (default is None)  
If None, do not try to decode the content of the files (e.g. for images or other non-text content). If not None, encoding to use to decode text files to Unicode if load\_content is True.

decode\_error : {'strict', 'ignore', 'replace'}, optional  
Instruction on what to do if a byte sequence is given to analyze that contains characters not of the given `encoding`. Passed as keyword argument 'errors' to bytes.decode.

random\_state : int, RandomState instance or None (default=0)  
Determines random number generation for dataset shuffling. Passes an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

Returns

-----

data : Bunch  
Dictionary-like object, the interesting attributes are: either data, the raw text data to learn, or 'filenames', the files holding it, 'target', the classification labels (integer indices), 'target\_names', the meaning of the labels, and 'DESCR', the full description of the dataset.

load\_iris(return\_X\_y=False)  
Load and return the iris dataset (classification).

The iris dataset is a classic and very easy multi-class classification dataset.

Classes	3
Samples per class	50
Samples total	150
Dimensionality	4
Features	real, positive

Read more in the :ref:`User Guide <iris\_dataset>`.

#### Parameters

-----

`return_X_y` : boolean, default=False.

If True, returns ``(data, target)`` instead of a Bunch object.

See

below for more information about the ``data`` and ``target`` objec

t.

.. versionadded:: 0.18

#### Returns

-----

`data` : Bunch

Dictionary-like object, the interesting attributes are:

'data', the data to learn, 'target', the classification label

s,

'target\_names', the meaning of the labels, 'feature\_names', th

e

meaning of the features, 'DESCR', the full description of  
the dataset, 'filename', the physical location of  
iris csv dataset (added in version `0.20`).

`(data, target)` : tuple if ``return\_X\_y`` is True

.. versionadded:: 0.18

#### Notes

-----

.. versionchanged:: 0.20

Fixed two wrong data points according to Fisher's paper.

The new version is the same as in R, but not as in the UCI  
Machine Learning Repository.

#### Examples

-----

Let's say you are interested in the samples 10, 25, and 50, and wa

nt to

know their class name.

```
>>> from sklearn.datasets import load_iris
>>> data = load_iris()
>>> data.target[[10, 25, 50]]
array([0, 0, 1])
>>> list(data.target_names)
['setosa', 'versicolor', 'virginica']
```

`load_linnerud(return_X_y=False)`

Load and return the linnerud dataset (multivariate regression).

```
=====
Samples total      20
Dimensionality     3 (for both data and target)
Features           integer
Targets            integer
=====
```

Read more in the :ref:`User Guide <linnerrud\_dataset>`.

#### Parameters

-----

return\_X\_y : boolean, default=False.  
If True, returns ``(data, target)`` instead of a Bunch object.  
See below for more information about the `data` and `target` o

bject.

.. versionadded:: 0.18

Returns

-----

data : Bunch

Dictionary-like object, the interesting attributes are: 'data'

and

'targets', the two multivariate datasets, with 'data' correspo

nding to

the exercise and 'targets' corresponding to the physiological measurements, as well as 'feature\_names' and 'target\_names'.

In addition, you will also have access to 'data\_filename',

the physical location of linnerud data csv dataset, and

'target\_filename', the physical location of

linnerud targets csv dataset (added in version `0.20`).

(data, target) : tuple if ``return\_X\_y`` is True

.. versionadded:: 0.18

load\_mlcomp(name\_or\_id, set\_='raw', mlcomp\_root=None, \*\*kwargs)

DEPRECATED: since the <http://mlcomp.org/> website will shut down in March 2017, the load\_mlcomp function was deprecated in version 0.19 and will be removed in 0.21.

Load a datasets as downloaded from <http://mlcomp.org>

Read more in the :ref:`User Guide <datasets>`.

Parameters

-----

name\_or\_id : int or str

The integer id or the string name metadata of the MLComp dataset to load

set\\_ : str, default='raw'

Select the portion to load: 'train', 'test' or 'raw'

mlcomp\_root : str, optional

The filesystem path to the root folder where MLComp dataset

ts

are stored, if mlcomp\_root is None, the MLCOMP\_DATASETS\_HO

ME

environment variable is looked up instead.

\*\*kwargs : domain specific kwargs to be passed to the dataset

loader.

Returns

-----

data : Bunch

Dictionary-like object, the interesting attributes are:

'filenames', the files holding the raw to learn, 'target',



the

classification labels (integer index), 'target\_names',  
the meaning of the labels, and 'DESCR', the full descripti

on of the

dataset.

Note on the lookup process: depending on the type of name\_or\_i  
d,

will choose between integer id lookup or metadata name lookup

by

looking at the unzipped archives and metadata file.

TODO: implement zip dataset loading too

load\_sample\_image(image\_name)

Load the numpy array of a single sample image

Read more in the :ref:`User Guide <sample\_images>`.

Parameters

-----

image\_name : {'china.jpg', 'flower.jpg'}

The name of the sample image loaded

Returns

-----

img : 3D array

The image as a numpy array: height x width x color

Examples

-----

```
>>> from sklearn.datasets import load_sample_image
>>> china = load_sample_image('china.jpg') # doctest: +SKIP
>>> china.dtype # doctest: +SKIP
dtype('uint8')
>>> china.shape # doctest: +SKIP
(427, 640, 3)
>>> flower = load_sample_image('flower.jpg') # doctest: +SKIP
>>> flower.dtype # doctest: +SKIP
dtype('uint8')
>>> flower.shape # doctest: +SKIP
(427, 640, 3)
```

load\_sample\_images()

Load sample images for image manipulation.

Loads both, ``china`` and ``flower``.

Read more in the :ref:`User Guide <sample\_images>`.

Returns

-----

data : Bunch

Dictionary-like object with the following attributes : 'image

s', the

two sample images, 'filenames', the file names for the images,

and

'DESCR' the full description of the dataset.

Examples

-----

To load the data and visualize the images:

```
>>> from sklearn.datasets import load_sample_images
>>> dataset = load_sample_images()      #doctest: +SKIP
>>> len(dataset.images)                 #doctest: +SKIP
2
>>> first_img_data = dataset.images[0] #doctest: +SKIP
>>> first_img_data.shape                 #doctest: +SKIP
(427, 640, 3)
>>> first_img_data.dtype                 #doctest: +SKIP
dtype('uint8')
```

```
load_svmlight_file(f, n_features=None, dtype=<class 'numpy.float64'>,
multilabel=False, zero_based='auto', query_id=False, offset=0, length=-1)
Load datasets in the svmlight / libsvm format into sparse CSR matrix
```

ix

This format is a text-based format, with one sample per line. It does not store zero valued features hence is suitable for sparse datasets.

The first element of each line can be used to store a target variable to predict.

This format is used as the default format for both svmlight and the libsvm command line programs.

Parsing a text based source can be expensive. When working on repeatedly on the same dataset, it is recommended to wrap this loader with `joblib.Memory.cache` to store a memmapped backup of the CSR results of the first call and benefit from the near instantaneous

ous

loading of memmapped structures for the subsequent calls.

In case the file contains a pairwise preference constraint (known as "qid" in the svmlight format) these are ignored unless the `query_id` parameter is set to `True`. These pairwise preference constraints can be used to constraint the combination of samples when using pairwise loss functions (as is the case in some learning to rank problems) so that only pairs with the same `query_id` value are considered.

This implementation is written in Cython and is reasonably fast. However, a faster API-compatible loader is also available at:

<https://github.com/mblondel/svmlight-loader>

Parameters

-----

`f` : {str, file-like, int}

t will

(Path to) a file to load. If a path ends in ".gz" or ".bz2", it

umed to

be uncompressed on the fly. If an integer is passed, it is assumed

be closed

to be a file descriptor. A file-like or file descriptor will not

be closed by this function. A file-like object must be opened in binary

mode.

`n_features : int or None`

The number of features to use. If None, it will be inferred. T

his

argument is useful to load several files that are subsets of a bigger sliced dataset: each subset might not have examples of every feature, hence the inferred shape might vary from one slice to another.

`n_features` is only required if `offset` or `length` are passed a non-default value.

`dtype : numpy data type, default np.float64`

Data type of dataset to be loaded. This will be the data type

of the

output numpy arrays `X` and `y`.

`multilabel : boolean, optional, default False`

Samples may have several labels each (see

<https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/multilabel.html>)

`zero_based : boolean or "auto", optional, default "auto"`

Whether column indices in `f` are zero-based (True) or one-based (False). If column indices are one-based, they are transformed

to

zero-based to match Python/NumPy conventions.

is from

If set to "auto", a heuristic check is applied to determine th

t they

the file contents. Both kinds of files occur "in the wild", bu

t they

are unfortunately not self-identifying. Using "auto" or True s

ould

always be safe when no `offset` or `length` is passed.

back

If `offset` or `length` are passed, the "auto" mode falls

ld

to `zero_based=True` to avoid having the heuristic check yie

inconsistent results on different segments of the file.

`query_id : boolean, default False`

If True, will return the `query_id` array for each file.

`offset : integer, optional, default 0`

Ignore the offset first bytes by seeking forward, then discarding the following bytes up until the next new line character.

`length : integer, optional, default -1`

If strictly positive, stop reading any new line of data once t

he

position in the file has reached the (offset + length) bytes t

hreshold.

Returns

-----

`X : scipy.sparse matrix of shape (n_samples, n_features)`

`y : ndarray of shape (n_samples,)`, or, in the multilabel a list of tuples of length `n_samples`.

query\_id : array of shape (n\_samples,)  
 query\_id for each sample. Only returned when query\_id is set to True.

See also

-----

load\_svmlight\_files: similar function for loading multiple files in this format, enforcing the same number of features/columns on all of them.

Examples

-----

To use joblib.Memory to cache the svmlight file::

```
from joblib import Memory
from sklearn.datasets import load_svmlight_file
mem = Memory("./mycache")
```

```
@mem.cache
def get_data():
    data = load_svmlight_file("mysvmlightfile")
    return data[0], data[1]
```

```
X, y = get_data()
```

```
load_svmlight_files(files, n_features=None, dtype=<class 'numpy.float64'>, multilabel=False, zero_based='auto', query_id=False, offset=0, length=-1)
```

Load dataset from multiple files in SVMlight format

This function is equivalent to mapping load\_svmlight\_file over a list of files, except that the results are concatenated into a single, flat list and the samples vectors are constrained to all have the same number of features.

In case the file contains a pairwise preference constraint (known as "qid" in the svmlight format) these are ignored unless the query\_id parameter is set to True. These pairwise preference constraints can be used to constraint the combination of samples when using pairwise loss functions (as is the case in some learning to rank problems) so that only pairs with the same query\_id value are considered.

Parameters

-----

files : iterable over {str, file-like, int}  
 (Paths of) files to load. If a path ends in ".gz" or ".bz2", it will be uncompressed on the fly. If an integer is passed, it is assumed to be a file descriptor. File-likes and file descriptors will not be closed by this function. File-like objects must be opened in binary mode.

`n_features` : int or None  
 The number of features to use. If None, it will be inferred from the maximum column index occurring in any of the files.

This can be set to a higher value than the actual number of features in any of the input files, but setting it to a lower value will cause an exception to be raised.

`dtype` : numpy data type, default `np.float64`  
 Data type of dataset to be loaded. This will be the data type of the output numpy arrays ```X``` and ```y```.

`multilabel` : boolean, optional  
 Samples may have several labels each (see <https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/multilabel.html>)

`zero_based` : boolean or "auto", optional  
 Whether column indices in `f` are zero-based (True) or one-based (False). If column indices are one-based, they are transformed to zero-based to match Python/NumPy conventions.  
 If set to "auto", a heuristic check is applied to determine the file contents. Both kinds of files occur "in the wild", but they are unfortunately not self-identifying. Using "auto" or True should always be safe when no offset or length is passed.  
 If offset or length are passed, the "auto" mode falls back to `zero_based=True` to avoid having the heuristic check yield inconsistent results on different segments of the file.

`query_id` : boolean, defaults to False  
 If True, will return the `query_id` array for each file.

`offset` : integer, optional, default 0  
 Ignore the offset first bytes by seeking forward, then discarding the following bytes up until the next new line character.

`length` : integer, optional, default -1  
 If strictly positive, stop reading any new line of data once the position in the file has reached the (offset + length) bytes threshold.

Returns  
 -----  
`[X1, y1, ..., Xn, yn]`  
 where each `(Xi, yi)` pair is the result from `load_svmlight_file(files[i])`.

If `query_id` is set to True, this will return instead `[X1, y1, q1, ..., Xn, yn, qn]` where `(Xi, yi, qi)` is the result from `load_svmlight_file(files[i])`

## Notes

-----

When fitting a model to a matrix `X_train` and evaluating it against a matrix `X_test`, it is essential that `X_train` and `X_test` have the same number of features (`X_train.shape[1] == X_test.shape[1]`). This may not be the case if you load the files individually with `load_svmlight_file`.

See also

-----

`load_svmlight_file``load_wine(return_X_y=False)`

Load and return the wine dataset (classification).

.. versionadded:: 0.18

The wine dataset is a classic and very easy multi-class classification dataset.

Classes	3
Samples per class	[59,71,48]
Samples total	178
Dimensionality	13
Features	real, positive

Read more in the :ref:`User Guide &lt;wine\_dataset&gt;`.

## Parameters

-----

`return_X_y` : boolean, default=False.

If True, returns ``(data, target)`` instead of a Bunch object.

See below for more information about the `data` and `target` objects.

## Returns

-----

`data` : Bunch

Dictionary-like object, the interesting attributes are: 'data', the data to learn, 'target', the classification labels, 'target\_names', the meaning of the labels, 'feature\_names', the meaning of the features, and 'DESCR', the full description of the dataset.

`(data, target)` : tuple if ``return\_X\_y`` is True

The copy of UCI ML Wine Data Set dataset is downloaded and modified to fit

standard format from:

<https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data>

## Examples

-----

ant to Let's say you are interested in the samples 10, 80, and 140, and w  
know their class name.

```
>>> from sklearn.datasets import load_wine
>>> data = load_wine()
>>> data.target[[10, 80, 140]]
array([0, 1, 2])
>>> list(data.target_names)
['class_0', 'class_1', 'class_2']
```

make\_biclusters(shape, n\_clusters, noise=0.0, minval=10, maxval=100, shuffle=True, random\_state=None)  
Generate an array with constant block diagonal structure for biclustering.

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

shape : iterable (n\_rows, n\_cols)  
The shape of the result.

n\_clusters : integer  
The number of biclusters.

noise : float, optional (default=0.0)  
The standard deviation of the gaussian noise.

minval : int, optional (default=10)  
Minimum value of a bicluster.

maxval : int, optional (default=100)  
Maximum value of a bicluster.

shuffle : boolean, optional (default=True)  
Shuffle the samples.

an int random\_state : int, RandomState instance or None (default)  
Determines random number generation for dataset creation. Pass  
for reproducible output across multiple function calls.  
See :term:`Glossary <random\_state>`.

#### Returns

-----

X : array of shape `shape`  
The generated array.

rows : array of shape (n\_clusters, X.shape[0],)  
The indicators for cluster membership of each row.

cols : array of shape (n\_clusters, X.shape[1],)  
The indicators for cluster membership of each column.

#### References

-----

.. [1] Dhillon, I. S. (2001, August). Co-clustering documents and words using bipartite spectral graph partitioning. In Proceedi

ngs

of the seventh ACM SIGKDD international conference on Knowledge

e

discovery and data mining (pp. 269-274). ACM.

See also

-----

make\_checkerboard

```
make_blobs(n_samples=100, n_features=2, centers=None, cluster_std=1.0,
center_box=(-10.0, 10.0), shuffle=True, random_state=None)
Generate isotropic Gaussian blobs for clustering.
```

Read more in the :ref:`User Guide &lt;sample\_generators&gt;`.

Parameters

-----

**n\_samples** : int or array-like, optional (default=100)  
 If int, it is the total number of points equally divided among clusters.  
 If array-like, each element of the sequence indicates the number of samples per cluster.

**n\_features** : int, optional (default=2)  
 The number of features for each sample.

**centers** : int or array of shape [n\_centers, n\_features], optional (default=None)  
 The number of centers to generate, or the fixed center locations.

ns.

rated.

If n\_samples is an int and centers is None, 3 centers are generated.

ples.

If n\_samples is array-like, centers must be either None or an array of length equal to the length of n\_samples.

**cluster\_std** : float or sequence of floats, optional (default=1.0)  
 The standard deviation of the clusters.

**center\_box** : pair of floats (min, max), optional (default=(-10.0, 10.0))

The bounding box for each cluster center when centers are generated at random.

**shuffle** : boolean, optional (default=True)  
 Shuffle the samples.

**random\_state** : int, RandomState instance or None (default)  
 Determines random number generation for dataset creation. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

Returns

-----

**X** : array of shape [n\_samples, n\_features]  
 The generated samples.

**y** : array of shape [n\_samples]  
 The integer labels for cluster membership of each sample.



## Examples

-----

```
>>> from sklearn.datasets.samples_generator import make_blobs
>>> X, y = make_blobs(n_samples=10, centers=3, n_features=2,
...                   random_state=0)
>>> print(X.shape)
(10, 2)
>>> y
array([0, 0, 1, 0, 2, 2, 2, 1, 1, 0])
>>> X, y = make_blobs(n_samples=[3, 3, 4], centers=None, n_feature
```

s=2,

```
...                   random_state=0)
>>> print(X.shape)
(10, 2)
>>> y
array([0, 1, 2, 0, 2, 2, 2, 1, 1, 0])
```

## See also

-----

make\_classification: a more intricate variant

```
make_checkerboard(shape, n_clusters, noise=0.0, minval=10, maxval=100,
shuffle=True, random_state=None)
```

Generate an array with block checkerboard structure for biclustering.

Read more in the :ref:`User Guide <sample\_generators>`.

## Parameters

-----

shape : iterable (n\_rows, n\_cols)  
The shape of the result.

n\_clusters : integer or iterable (n\_row\_clusters, n\_column\_clusters)  
The number of row and column clusters.

noise : float, optional (default=0.0)  
The standard deviation of the gaussian noise.

minval : int, optional (default=10)  
Minimum value of a bicluster.

maxval : int, optional (default=100)  
Maximum value of a bicluster.

shuffle : boolean, optional (default=True)  
Shuffle the samples.

random\_state : int, RandomState instance or None (default)  
Determines random number generation for dataset creation. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

## Returns

-----

X : array of shape `shape`  
The generated array.

rows : array of shape (n\_clusters, X.shape[0],)

The indicators for cluster membership of each row.

cols : array of shape (n\_clusters, X.shape[1],)  
The indicators for cluster membership of each column.

## References

-----

.. [1] Kluger, Y., Basri, R., Chang, J. T., & Gerstein, M. (2003).  
Spectral biclustering of microarray data: coclustering genes  
and conditions. Genome research, 13(4), 703-716.

See also

-----

make\_biclusters

make\_circles(n\_samples=100, shuffle=True, noise=None, random\_state=None, factor=0.8)

Make a large circle containing a smaller circle in 2d.

A simple toy dataset to visualize clustering and classification algorithms.

Read more in the :ref:`User Guide <sample\_generators>`.

## Parameters

-----

n\_samples : int, optional (default=100)

The total number of points generated. If odd, the inner circle  
will have one point more than the outer circle.

shuffle : bool, optional (default=True)

Whether to shuffle the samples.

noise : double or None (default=None)

Standard deviation of Gaussian noise added to the data.

random\_state : int, RandomState instance or None (default)

Determines random number generation for dataset shuffling and

noise.

Pass an int for reproducible output across multiple function c

alls.

See :term:`Glossary <random\_state>`.

factor : 0 < double < 1 (default=.8)

Scale factor between inner and outer circle.

## Returns

-----

X : array of shape [n\_samples, 2]

The generated samples.

y : array of shape [n\_samples]

The integer labels (0 or 1) for class membership of each sampl  
e.

make\_classification(n\_samples=100, n\_features=20, n\_informative=2, n\_r  
edundant=2, n\_repeated=0, n\_classes=2, n\_clusters\_per\_class=2, weights=None, flip\_y=0.01, class\_sep=1.0, hypercube=True, shift=0.0, scale=1.0, shuff

```
le=True, random_state=None)
```

```
    Generate a random n-class classification problem.
```

```
    This initially creates clusters of points normally distributed (std=1)
```

```
    about vertices of an ``n_informative``-dimensional hypercube with
    sides of length ``2*class_sep`` and assigns an equal number of clusters to
    each
```

```
    class. It introduces interdependence between these features and adds
    various types of further noise to the data.
```

```
    Without shuffling, ``X`` horizontally stacks features in the following
```

```
    order: the primary ``n_informative`` features, followed by ``n_redundant``
```

```
    linear combinations of the informative features, followed by ``n_repeated``
```

```
    duplicates, drawn randomly with replacement from the informative and
```

```
    redundant features. The remaining features are filled with random noise.
```

```
    Thus, without shuffling, all useful features are contained in the columns
```

```
    ``X[:, :n_informative + n_redundant + n_repeated]``.
```

```
    Read more in the :ref:`User Guide <sample_generators>`.
```

```
Parameters
```

```
-----
```

```
n_samples : int, optional (default=100)
```

```
    The number of samples.
```

```
n_features : int, optional (default=20)
```

```
    The total number of features. These comprise ``n_informative`` informative
    features, ``n_redundant`` redundant features, ``n_repeated`` duplicated
    features and ``n_features-n_informative-n_redundant-n_repeated`` useless fe
```

```
atures
```

```
    drawn at random.
```

```
n_informative : int, optional (default=2)
```

```
    The number of informative features. Each class is composed of a number
```

```
    of gaussian clusters each located around the vertices of a hypercube
```

```
    in a subspace of dimension ``n_informative``. For each cluster,
```

```
    informative features are drawn independently from  $N(0, 1)$  and then
```

```
    randomly linearly combined within each cluster in order to add covariance.
    The clusters are then placed on the vertices of the
```

```
hypercube.
```

```
n_redundant : int, optional (default=2)
```

```
    The number of redundant features. These features are generated as
```

```
    random linear combinations of the informative features.
```

`n_repeated : int, optional (default=0)`  
 The number of duplicated features, drawn randomly from the informative and the redundant features.

`n_classes : int, optional (default=2)`  
 The number of classes (or labels) of the classification problem.

`n_clusters_per_class : int, optional (default=2)`  
 The number of clusters per class.

`weights : list of floats or None (default=None)`  
 The proportions of samples assigned to each class. If None, then classes are balanced. Note that if `len(weights) == n_classes - 1`, then the last class weight is automatically inferred. More than `n_samples` samples may be returned if the sum of `weights` exceeds 1.

`flip_y : float, optional (default=0.01)`  
 The fraction of samples whose class are randomly exchanged. Larger values introduce noise in the labels and make the classification task harder.

`class_sep : float, optional (default=1.0)`  
 The factor multiplying the hypercube size. Larger values spread out the clusters/classes and make the classification task easier.

`hypercube : boolean, optional (default=True)`  
 If True, the clusters are put on the vertices of a hypercube. If False, the clusters are put on the vertices of a random polytope.

`shift : float, array of shape [n_features] or None, optional (default=0.0)`  
 Shift features by the specified value. If None, then features are shifted by a random value drawn in `[-class_sep, class_sep]`.

`scale : float, array of shape [n_features] or None, optional (default=1.0)`  
 Multiply features by the specified value. If None, then features are scaled by a random value drawn in `[1, 100]`. Note that scaling happens after shifting.

`shuffle : boolean, optional (default=True)`  
 Shuffle the samples and the features.

`random_state : int, RandomState instance or None (default)`  
 Determines random number generation for dataset creation. Pass an int for reproducible output across multiple function calls.

See :term:`Glossary <random\_state>`.

#### Returns

-----

`X` : array of shape `[n_samples, n_features]`  
The generated samples.

`y` : array of shape `[n_samples]`  
The integer labels for class membership of each sample.

#### Notes

-----

The algorithm is adapted from Guyon [1] and was designed to generate the "Madelon" dataset.

#### References

-----

.. [1] I. Guyon, "Design of experiments for the NIPS 2003 variable selection benchmark", 2003.

#### See also

-----

`make_blobs`: simplified variant

`make_multilabel_classification`: unrelated generator for multilabel

tasks

`make_friedman1(n_samples=100, n_features=10, noise=0.0, random_state=None)`

Generate the "Friedman #1" regression problem

This dataset is described in Friedman [1] and Breiman [2].

Inputs ``X`` are independent features uniformly distributed on the interval

`[0, 1]`. The output ``y`` is created according to the formula::

$$y(X) = 10 * \sin(\pi * X[:, 0] * X[:, 1]) + 20 * (X[:, 2] - 0.5) ** 2 + 10 * X[:, 3] + 5 * X[:, 4] + \text{noise} * N(0, 1).$$

Out of the ``n_features`` features, only 5 are actually used to compute ``y``. The remaining features are independent of ``y``.

The number of features has to be `>= 5`.

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : int, optional (default=100)  
The number of samples.

`n_features` : int, optional (default=10)  
The number of features. Should be at least 5.

`noise` : float, optional (default=0.0)  
The standard deviation of the gaussian noise applied to the output.

put.

`random_state` : int, RandomState instance or None (default)

`int` Determines random number generation for dataset noise. Pass an `int` for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

#### Returns

-----

`X` : array of shape `[n_samples, n_features]`  
The input samples.

`y` : array of shape `[n_samples]`  
The output values.

#### References

-----

.. [1] J. Friedman, "Multivariate adaptive regression splines", *The Annals of Statistics* 19 (1), pages 1-67, 1991.

.. [2] L. Breiman, "Bagging predictors", *Machine Learning* 24, pages 123-140, 1996.

`make_friedman2(n_samples=100, noise=0.0, random_state=None)`  
Generate the "Friedman #2" regression problem

This dataset is described in Friedman [1] and Breiman [2].

Inputs ``X`` are 4 independent features uniformly distributed on the intervals::

```
0 <= X[:, 0] <= 100,
40 * pi <= X[:, 1] <= 560 * pi,
0 <= X[:, 2] <= 1,
1 <= X[:, 3] <= 11.
```

The output ``y`` is created according to the formula::

$$y(X) = (X[:, 0] ** 2 + (X[:, 1] * X[:, 2] - 1 / (X[:, 1] * X[:, 3])) ** 2) ** 0.5 + \text{noise} * N(0, 1).$$

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : `int`, optional (default=100)  
The number of samples.

`noise` : `float`, optional (default=0.0)  
The standard deviation of the gaussian noise applied to the output.

`random_state` : `int`, `RandomState` instance or `None` (default)  
Determines random number generation for dataset noise. Pass an `int` for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

#### Returns

-----

`X` : array of shape `[n_samples, 4]`  
The input samples.

`y` : array of shape `[n_samples]`  
The output values.

#### References

-----

.. [1] J. Friedman, "Multivariate adaptive regression splines", *The Annals of Statistics* 19 (1), pages 1-67, 1991.

.. [2] L. Breiman, "Bagging predictors", *Machine Learning* 24, pages 123-140, 1996.

`make_friedman3(n_samples=100, noise=0.0, random_state=None)`  
Generate the "Friedman #3" regression problem

This dataset is described in Friedman [1] and Breiman [2].

Inputs ``X`` are 4 independent features uniformly distributed on the intervals::

```
0 <= X[:, 0] <= 100,
40 * pi <= X[:, 1] <= 560 * pi,
0 <= X[:, 2] <= 1,
1 <= X[:, 3] <= 11.
```

The output ``y`` is created according to the formula::

$$y(X) = \arctan((X[:, 1] * X[:, 2] - 1 / (X[:, 1] * X[:, 3])) / X[:, 0]) + \text{noise} * N(0, 1).$$

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : int, optional (default=100)  
The number of samples.

`noise` : float, optional (default=0.0)  
The standard deviation of the gaussian noise applied to the output.

`random_state` : int, RandomState instance or None (default)  
Determines random number generation for dataset noise. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

#### Returns

-----

`X` : array of shape `[n_samples, 4]`  
The input samples.

`y` : array of shape `[n_samples]`  
The output values.

#### References

-----

.. [1] J. Friedman, "Multivariate adaptive regression splines", *The Annals of Statistics* 19 (1), pages 1-67, 1991.

.. [2] L. Breiman, "Bagging predictors", Machine Learning 24, pages 123-140, 1996.

```
make_gaussian_quantiles(mean=None, cov=1.0, n_samples=100, n_features=
2, n_classes=3, shuffle=True, random_state=None)
```

Generate isotropic Gaussian and label samples by quantile

This classification dataset is constructed by taking a multi-dimensional standard normal distribution and defining classes separated by nested concentric multi-dimensional spheres such that roughly equal numbers of samples are in each class (quantiles of the  $\chi^2$  distribution).

Read more in the :ref:`User Guide <sample\_generators>`.

Parameters

-----

mean : array of shape [n\_features], optional (default=None)

The mean of the multi-dimensional normal distribution.

If None then use the origin (0, 0, ...).

cov : float, optional (default=1.)

The covariance matrix will be this value times the unit matrix.

x. This

dataset only produces symmetric normal distributions.

n\_samples : int, optional (default=100)

The total number of points equally divided among classes.

n\_features : int, optional (default=2)

The number of features for each sample.

n\_classes : int, optional (default=3)

The number of classes

shuffle : boolean, optional (default=True)

Shuffle the samples.

random\_state : int, RandomState instance or None (default)

Determines random number generation for dataset creation. Pass

an int

for reproducible output across multiple function calls.

See :term:`Glossary <random\_state>`.

Returns

-----

X : array of shape [n\_samples, n\_features]

The generated samples.

y : array of shape [n\_samples]

The integer labels for quantile membership of each sample.

Notes

-----

The dataset is from Zhu et al [1].

References



-----  
 .. [1] J. Zhu, H. Zou, S. Rosset, T. Hastie, "Multi-class AdaBoos  
 t", 2009.

make\_hastie\_10\_2(n\_samples=12000, random\_state=None)  
 Generates data for binary classification used in  
 Hastie et al. 2009, Example 10.2.

The ten features are standard independent Gaussian and  
 the target ``y`` is defined by::

$y[i] = 1$  if  $\text{np.sum}(X[i] ** 2) > 9.34$  else  $-1$

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

n\_samples : int, optional (default=12000)

The number of samples.

random\_state : int, RandomState instance or None (default)

Determines random number generation for dataset creation. Pass

an int

for reproducible output across multiple function calls.

See :term:`Glossary <random\_state>`.

#### Returns

-----

X : array of shape [n\_samples, 10]

The input samples.

y : array of shape [n\_samples]

The output values.

#### References

-----

.. [1] T. Hastie, R. Tibshirani and J. Friedman, "Elements of Stat  
 istical  
 Learning Ed. 2", Springer, 2009.

#### See also

-----

make\_gaussian\_quantiles: a generalization of this dataset approach

make\_low\_rank\_matrix(n\_samples=100, n\_features=100, effective\_rank=10,  
 tail\_strength=0.5, random\_state=None)

Generate a mostly low rank matrix with bell-shaped singular values

Most of the variance can be explained by a bell-shaped curve of wi  
 dth

effective\_rank: the low rank part of the singular values profile i  
 s::

$(1 - \text{tail\_strength}) * \exp(-1.0 * (i / \text{effective\_rank}) ** 2)$

The remaining singular values' tail is fat, decreasing as::

$\text{tail\_strength} * \exp(-0.1 * i / \text{effective\_rank}).$

The low rank part of the profile can be considered the structured  
 signal part of the data while the tail can be considered the noisy

part of the data that cannot be summarized by a low number of linear components (singular vectors).

This kind of singular profiles is often seen in practice, for instance:

- gray level pictures of faces
- TF-IDF vectors of text documents crawled from the web

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : int, optional (default=100)  
The number of samples.

`n_features` : int, optional (default=100)  
The number of features.

`effective_rank` : int, optional (default=10)  
The approximate number of singular vectors required to explain most of the data by linear combinations.

`tail_strength` : float between 0.0 and 1.0, optional (default=0.5)  
The relative importance of the fat noisy tail of the singular values profile.

`random_state` : int, RandomState instance or None (default)  
Determines random number generation for dataset creation. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

#### Returns

-----

`X` : array of shape [n\_samples, n\_features]  
The matrix.

`make_moons(n_samples=100, shuffle=True, noise=None, random_state=None)`  
Make two interleaving half circles

A simple toy dataset to visualize clustering and classification algorithms. Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : int, optional (default=100)  
The total number of points generated.

`shuffle` : bool, optional (default=True)  
Whether to shuffle the samples.

`noise` : double or None (default=None)  
Standard deviation of Gaussian noise added to the data.

`random_state` : int, RandomState instance or None (default)  
Determines random number generation for dataset shuffling and noise.

Pass an int for reproducible output across multiple function calls.

See :term:`Glossary <random\_state>`.

#### Returns

-----

`X` : array of shape `[n_samples, 2]`  
The generated samples.

`y` : array of shape `[n_samples]`  
The integer labels (0 or 1) for class membership of each sample.

`make_multilabel_classification(n_samples=100, n_features=20, n_classes=5, n_labels=2, length=50, allow_unlabeled=True, sparse=False, return_indicator='dense', return_distributions=False, random_state=None)`

Generate a random multilabel classification problem.

For each sample, the generative process is:

- pick the number of labels:  $n \sim \text{Poisson}(n\_labels)$
- $n$  times, choose a class  $c$ :  $c \sim \text{Multinomial}(\theta_c)$
- pick the document length:  $k \sim \text{Poisson}(\text{length})$
- $k$  times, choose a word:  $w \sim \text{Multinomial}(\theta_c)$

In the above process, rejection sampling is used to make sure that  $n$  is never zero or more than `n_classes`, and that the document length is never zero. Likewise, we reject classes which have already been chosen.

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : int, optional (default=100)  
The number of samples.

`n_features` : int, optional (default=20)  
The total number of features.

`n_classes` : int, optional (default=5)  
The number of classes of the classification problem.

`n_labels` : int, optional (default=2)  
The average number of labels per instance. More precisely, the number of labels per sample is drawn from a Poisson distribution with `n_labels` as its expected value, but samples are bounded (using rejection sampling) by `n_classes`, and must be nonzero if `allow_unlabeled` is False.

`length` : int, optional (default=50)  
The sum of the features (number of words if documents) is drawn from a Poisson distribution with this expected value.

`allow_unlabeled` : bool, optional (default=True)  
If `True`, some instances might not belong to any class.

`sparse` : bool, optional (default=False)

If ``True``, return a sparse feature matrix

.. versionadded:: 0.17  
parameter to allow *sparse* output.

return\_indicator : 'dense' (default) | 'sparse' | False  
If ``dense`` return ``Y`` in the dense binary indicator format.  
If  
t. If  
t. ``'sparse'`` return ``Y`` in the sparse binary indicator format.  
t. ``False`` returns a list of lists of labels.

return\_distributions : bool, optional (default=False)  
If ``True``, return the prior class probability and conditional  
1 probabilities of features given classes, from which the data was  
as drawn.

random\_state : int, RandomState instance or None (default)  
Determines random number generation for dataset creation. Pass  
an int  
for reproducible output across multiple function calls.  
See :term:`Glossary <random\_state>`.

Returns

-----

X : array of shape [n\_samples, n\_features]  
The generated samples.

Y : array or sparse CSR matrix of shape [n\_samples, n\_classes]  
The label sets.

p\_c : array, shape [n\_classes]  
The probability of each class being drawn. Only returned if  
``return\_distributions=True``.

p\_w\_c : array, shape [n\_features, n\_classes]  
The probability of each feature being drawn given each class.  
Only returned if ``return\_distributions=True``.

make\_regression(n\_samples=100, n\_features=100, n\_informative=10, n\_targets=1, bias=0.0, effective\_rank=None, tail\_strength=0.5, noise=0.0, shuffle=True, coef=False, random\_state=None)  
Generate a random regression problem.

The input set can either be well conditioned (by default) or have  
a low  
rank-fat tail singular profile. See :func:`make\_low\_rank\_matrix` for  
or more details.

The output is generated by applying a (potentially biased) random  
linear  
regression model with `n\_informative` nonzero regressors to the previously  
generated input and some gaussian centered noise with some adjustable  
ble scale.

Read more in the :ref:`User Guide <sample\_generators>`.

## Parameters

-----

`n_samples : int, optional (default=100)`

The number of samples.

`n_features : int, optional (default=100)`

The number of features.

`n_informative : int, optional (default=10)`

es used  
to build the linear model used to generate the output.

`n_targets : int, optional (default=1)`

output  
The number of regression targets, i.e., the dimension of the y  
vector associated with a sample. By default, the output is a scalar.

`bias : float, optional (default=0.0)`

The bias term in the underlying linear model.

`effective_rank : int or None, optional (default=None)`

lain most  
of  
rodude  
ith  
if not None:  
The approximate number of singular vectors required to exp  
of the input data by linear combinations. Using this kind  
singular spectrum in the input allows the generator to rep  
the correlations often observed in practice.  
if None:  
The input set is well conditioned, centered and gaussian w  
unit variance.

`tail_strength : float between 0.0 and 1.0, optional (default=0.5)`

values  
The relative importance of the fat noisy tail of the singular  
profile if `effective\_rank` is not None.

`noise : float, optional (default=0.0)`

tput.  
The standard deviation of the gaussian noise applied to the ou

`shuffle : boolean, optional (default=True)`

Shuffle the samples and the features.

`coef : boolean, optional (default=False)`

eturned.  
If True, the coefficients of the underlying linear model are r

`random_state : int, RandomState instance or None (default)`

an int  
Determines random number generation for dataset creation. Pass  
for reproducible output across multiple function calls.  
See :term:`Glossary <random\_state>`.

## Returns

-----

`X` : array of shape `[n_samples, n_features]`  
The input samples.

`y` : array of shape `[n_samples]` or `[n_samples, n_targets]`  
The output values.

`coef` : array of shape `[n_features]` or `[n_features, n_targets]`, optional  
The coefficient of the underlying linear model. It is returned only if `coef` is `True`.

`make_s_curve(n_samples=100, noise=0.0, random_state=None)`  
Generate an S curve dataset.

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : int, optional (default=100)  
The number of sample points on the S curve.

`noise` : float, optional (default=0.0)  
The standard deviation of the gaussian noise.

`random_state` : int, RandomState instance or None (default)  
Determines random number generation for dataset creation. Pass an int for reproducible output across multiple function calls. See :term:`Glossary <random\_state>`.

#### Returns

-----

`X` : array of shape `[n_samples, 3]`  
The points.

`t` : array of shape `[n_samples]`  
The univariate position of the sample according to the main dimension of the points in the manifold.

`make_sparse_coded_signal(n_samples, n_components, n_features, n_nonzero_coefs, random_state=None)`  
Generate a signal as a sparse combination of dictionary elements.

Returns a matrix  $Y = DX$ , such as  $D$  is `(n_features, n_components)`,  $X$  is `(n_components, n_samples)` and each column of  $X$  has exactly `n_nonzero_coefs` non-zero elements.

Read more in the :ref:`User Guide <sample\_generators>`.

#### Parameters

-----

`n_samples` : int  
number of samples to generate

`n_components` : int,  
number of components in the dictionary

`n_features` : int  
number of features of the dataset to generate

`n_nonzero_coefs : int`  
 number of active (non-zero) coefficients in each sample

`random_state : int, RandomState instance or None (default)`  
 Determines random number generation for dataset creation. Pass  
 an int  
 for reproducible output across multiple function calls.  
 See :term:`Glossary <random\_state>`.

Returns

-----

`data : array of shape [n_features, n_samples]`  
 The encoded signal (Y).

`dictionary : array of shape [n_features, n_components]`  
 The dictionary with normalized components (D).

`code : array of shape [n_components, n_samples]`  
 The sparse code such that each column of this matrix has exact  
 ly  
`n_nonzero_coefs` non-zero items (X).

`make_sparse_spd_matrix(dim=1, alpha=0.95, norm_diag=False, smallest_coef=0.1, largest_coef=0.9, random_state=None)`  
 Generate a sparse symmetric definite positive matrix.

Read more in the :ref:`User Guide <sample\_generators>`.

Parameters

-----

`dim : integer, optional (default=1)`  
 The size of the random matrix to generate.

`alpha : float between 0 and 1, optional (default=0.95)`  
 The probability that a coefficient is zero (see notes). Larger  
 values  
 enforce more sparsity.

`norm_diag : boolean, optional (default=False)`  
 Whether to normalize the output matrix to make the leading dia  
 gonal  
 elements all 1

`smallest_coef : float between 0 and 1, optional (default=0.1)`  
 The value of the smallest coefficient.

`largest_coef : float between 0 and 1, optional (default=0.9)`  
 The value of the largest coefficient.

`random_state : int, RandomState instance or None (default)`  
 Determines random number generation for dataset creation. Pass  
 an int  
 for reproducible output across multiple function calls.  
 See :term:`Glossary <random\_state>`.

Returns

-----

`prec : sparse matrix of shape (dim, dim)`  
 The generated matrix.

## Notes

-----

rix.  
f  
The sparsity is actually imposed on the cholesky factor of the matrix.  
Thus alpha does not translate directly into the filling fraction of the matrix itself.

See also

-----

make\_spd\_matrix

make\_sparse\_uncorrelated(n\_samples=100, n\_features=10, random\_state=None)  
gn  
Generate a random regression problem with sparse uncorrelated design

This dataset is described in Celeux et al [1]. as::

$$X \sim N(0, 1)$$

$$y(X) = X[:, 0] + 2 * X[:, 1] - 2 * X[:, 2] - 1.5 * X[:, 3]$$

are  
Only the first 4 features are informative. The remaining features are useless.

Read more in the :ref:`User Guide &lt;sample\_generators&gt;`.

## Parameters

-----

n\_samples : int, optional (default=100)  
The number of samples.

n\_features : int, optional (default=10)  
The number of features.

random\_state : int, RandomState instance or None (default)  
Determines random number generation for dataset creation. Pass an int  
for reproducible output across multiple function calls.  
See :term:`Glossary <random\_state>`.

## Returns

-----

X : array of shape [n\_samples, n\_features]  
The input samples.

y : array of shape [n\_samples]  
The output values.

## References

-----

.. [1] G. Celeux, M. El Anbari, J.-M. Marin, C. P. Robert,  
entist  
"Regularization in regression: comparing Bayesian and frequentist methods in a poorly informative situation", 2009.

make\_spd\_matrix(n\_dim, random\_state=None)  
Generate a random symmetric, positive-definite matrix.

Read more in the :ref:`User Guide &lt;sample\_generators&gt;`.



## Parameters

-----

`n_dim : int`

The matrix dimension.

`random_state : int, RandomState instance or None (default)`

Determines random number generation for dataset creation. Pass

an int

for reproducible output across multiple function calls.

See :term:`Glossary &lt;random\_state&gt;`.

## Returns

-----

`X : array of shape [n_dim, n_dim]`

The random symmetric, positive-definite matrix.

## See also

-----

`make_sparse_spd_matrix``make_swiss_roll(n_samples=100, noise=0.0, random_state=None)`

Generate a swiss roll dataset.

Read more in the :ref:`User Guide &lt;sample\_generators&gt;`.

## Parameters

-----

`n_samples : int, optional (default=100)`

The number of sample points on the S curve.

`noise : float, optional (default=0.0)`

The standard deviation of the gaussian noise.

`random_state : int, RandomState instance or None (default)`

Determines random number generation for dataset creation. Pass

an int

for reproducible output across multiple function calls.

See :term:`Glossary &lt;random\_state&gt;`.

## Returns

-----

`X : array of shape [n_samples, 3]`

The points.

`t : array of shape [n_samples]`

The univariate position of the sample according to the main di

mension

of the points in the manifold.

## Notes

-----

The algorithm is from Marsland [1].

## References

-----

.. [1] S. Marsland, "Machine Learning: An Algorithmic Perspectiv

e",

Chapter 10, 2009.

<http://seat.massey.ac.nz/personal/s.r.marsland/Code/10/11e>.

py

```
mldata_filename(dataname)
DEPRECATED: mldata_filename was deprecated in version 0.20 and will
be removed in version 0.22
```

Convert a raw name for a data set in a mldata.org filename.

```
.. deprecated:: 0.20
    Will be removed in version 0.22
```

Parameters

-----

```
dataname : str
    Name of dataset
```

Returns

-----

```
fname : str
    The converted dataname.
```

DATA

```
__all__ = ['clear_data_home', 'dump_svmlight_file', 'fetch_20newsgroups',
p...
```

FILE

```
/home/a/lectenv/lib/python3.6/site-packages/sklearn/datasets/__init__.
py
```

## Scripts

- Back to Python as a glue.
- Take us outside Notebooks.

In [ ]:

```
#!/usr/bin/env python3

# Why do we want the comment above?
# Tell us the interpreter.

if __name__ == "__main__":
    print("Running as main.")
else:
    print("I was imported!")
```

After lecture: During demo in class, I made the file executable by `chmod +x foo.py` (as this was the name of my script).

- Several useful modules to check out. For instance `sys`, `subprocess` (OS commands).

# Pointers to noteworthy data science packages

- [matplotlib](https://matplotlib.org/) (<https://matplotlib.org/>) (home page with tutorials)
  - [scikit-learn](https://scikit-learn.org/stable/index.html) (<https://scikit-learn.org/stable/index.html>) (import as sklearn)
  - [numpy](http://www.numpy.org/) (<http://www.numpy.org/>). Arrays!
  - [scipy](https://www.scipy.org/) (<https://www.scipy.org/>) Note: sparse matrices.
  - [pandas](https://pandas.pydata.org/) (<https://pandas.pydata.org/>)
- 
- Entire courses.

## Wrapping up